

Fundamentals of Designing a Data Warehouse

Sensible techniques for developing a data warehousing environment which is relevant, agile, and extensible

Melissa Coates
BI Architect, SentryOne
sentryone.com



Blog: sqlchick.com
Twitter: [@sqlchick](https://twitter.com/sqlchick)

Fundamentals of Designing a Data Warehouse

Agenda

1. Overview of the Need for Data Warehousing
2. DW Design Principles
3. Dimension Design
4. Fact Design
5. When to Use Columnstore or Partitioning
6. DW Tips
7. SSDT 'Database Project' Tips
8. Planning Future Growth of the DW

*All syntax shown is from
SQL Server 2016.*

*Screen shots are from
SQL Server Data Tools in
Visual Studio 2015.*

Fundamentals of Designing a Data Warehouse

Out of Scope

- ✓ ETL patterns and techniques
- ✓ Source control
- ✓ Deployment practices
- ✓ Master data management
- ✓ Data quality techniques
- ✓ Semantic layer, OLAP, cubes
- ✓ Front-end reporting
- ✓ Security
- ✓ Tuning & monitoring
- ✓ Automation techniques

Overview of the Need for Data Warehousing

First Let's Get This Straight...

Data Warehousing is not dead!

Data warehousing can be “uncool” but it doesn't have to be if you adopt modern data warehousing concepts & technologies such as:

- ✓ Data lake
- ✓ Hadoop
- ✓ Real-time
- ✓ Large data volume
- ✓ Data virtualization
- ✓ Hybrid & cloud
- ✓ Automation
- ✓ Bimodal environments

Transaction System vs. Data Warehouse

OLTP

Goal:

- ✓ Operational transactions
- ✓ "Writes"

Scope:

One database system

Example Objectives:

- ✓ Process a customer order
- ✓ Generate an invoice

Data Warehouse

Goal:

- ✓ Informational and analytical
- ✓ "Reads"

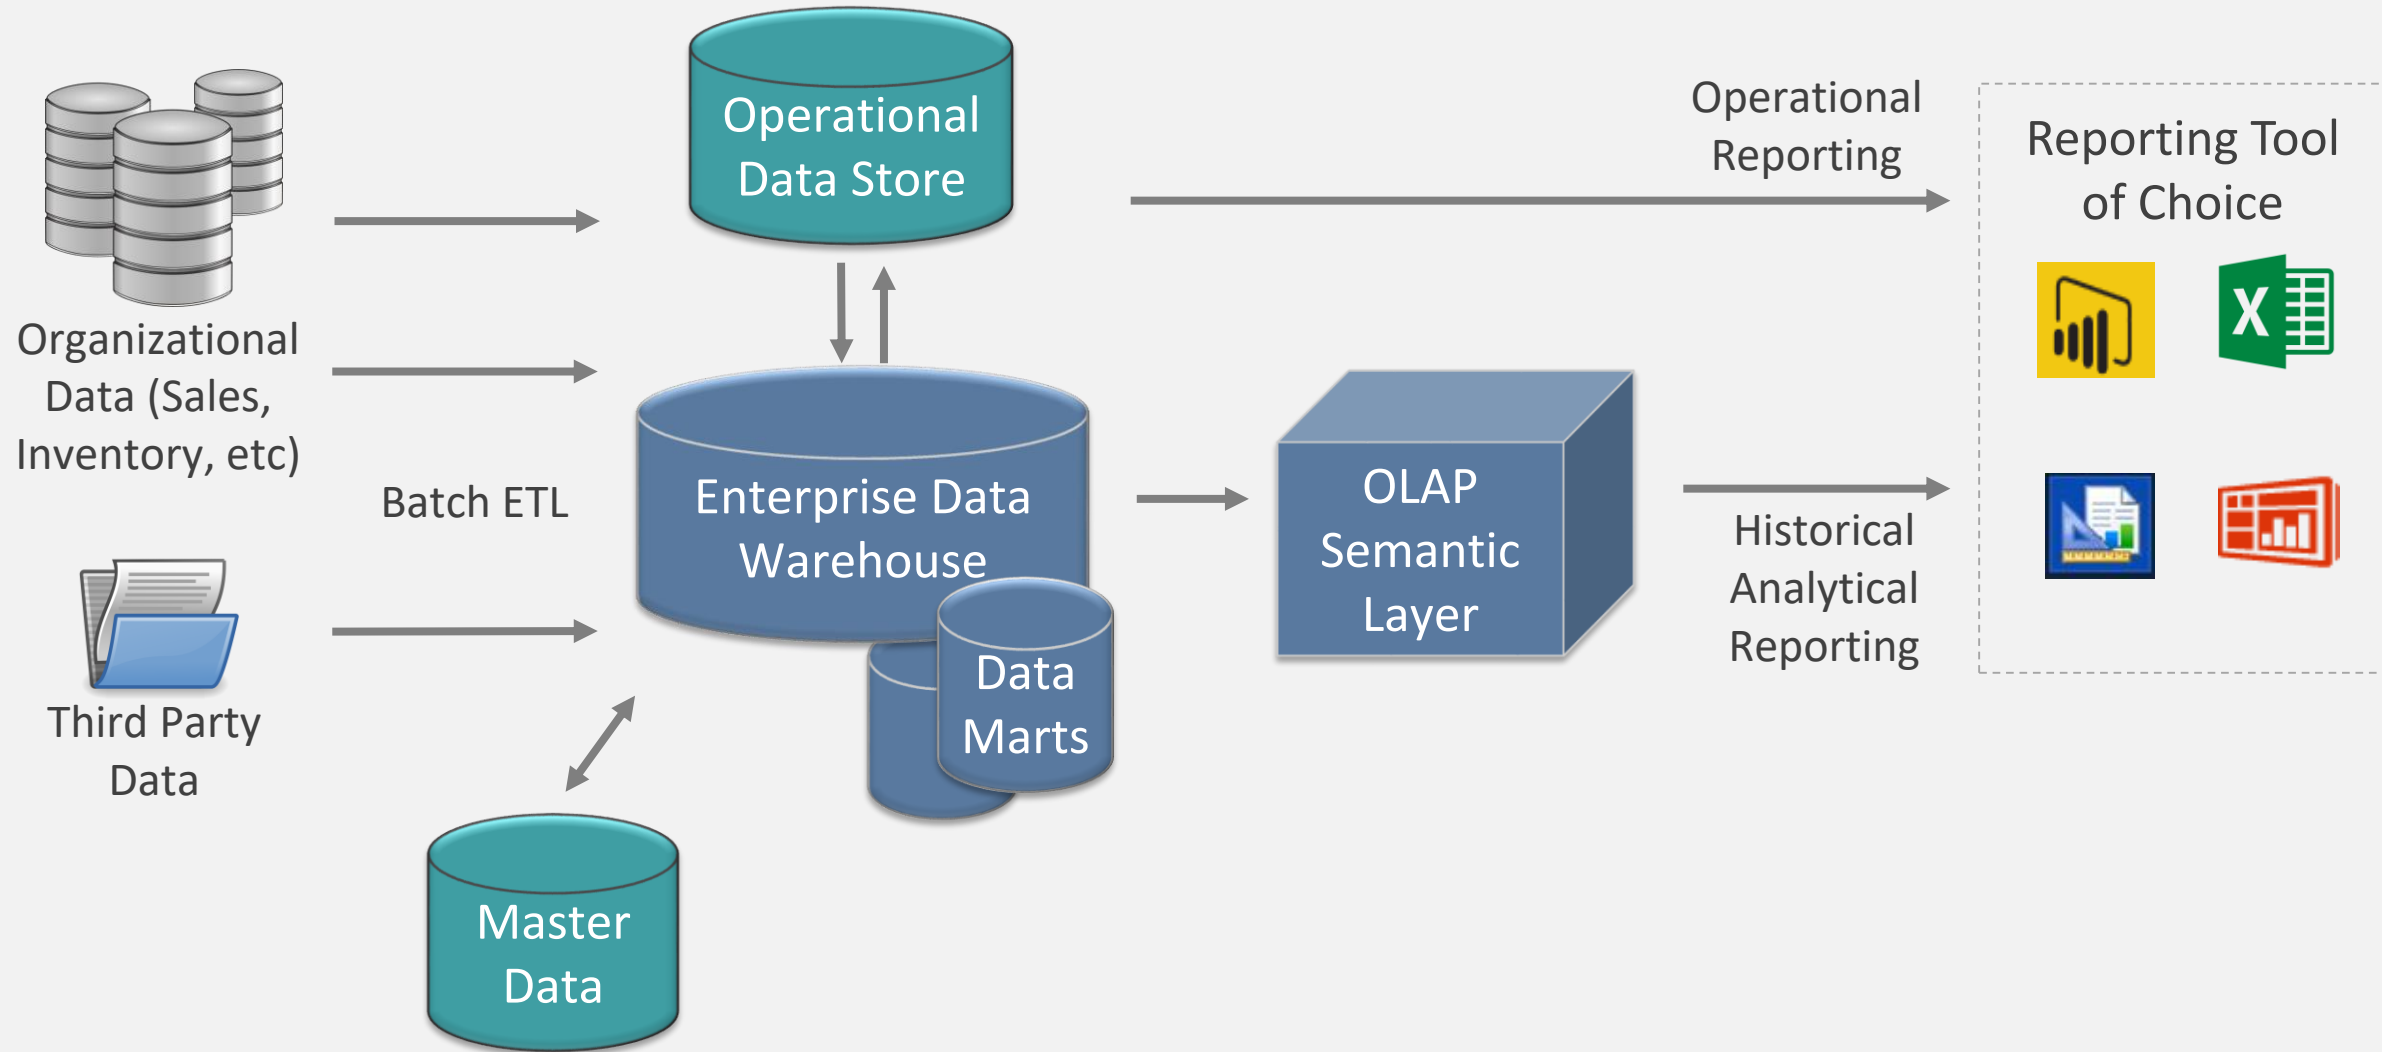
Scope:

Integrate data from multiple systems

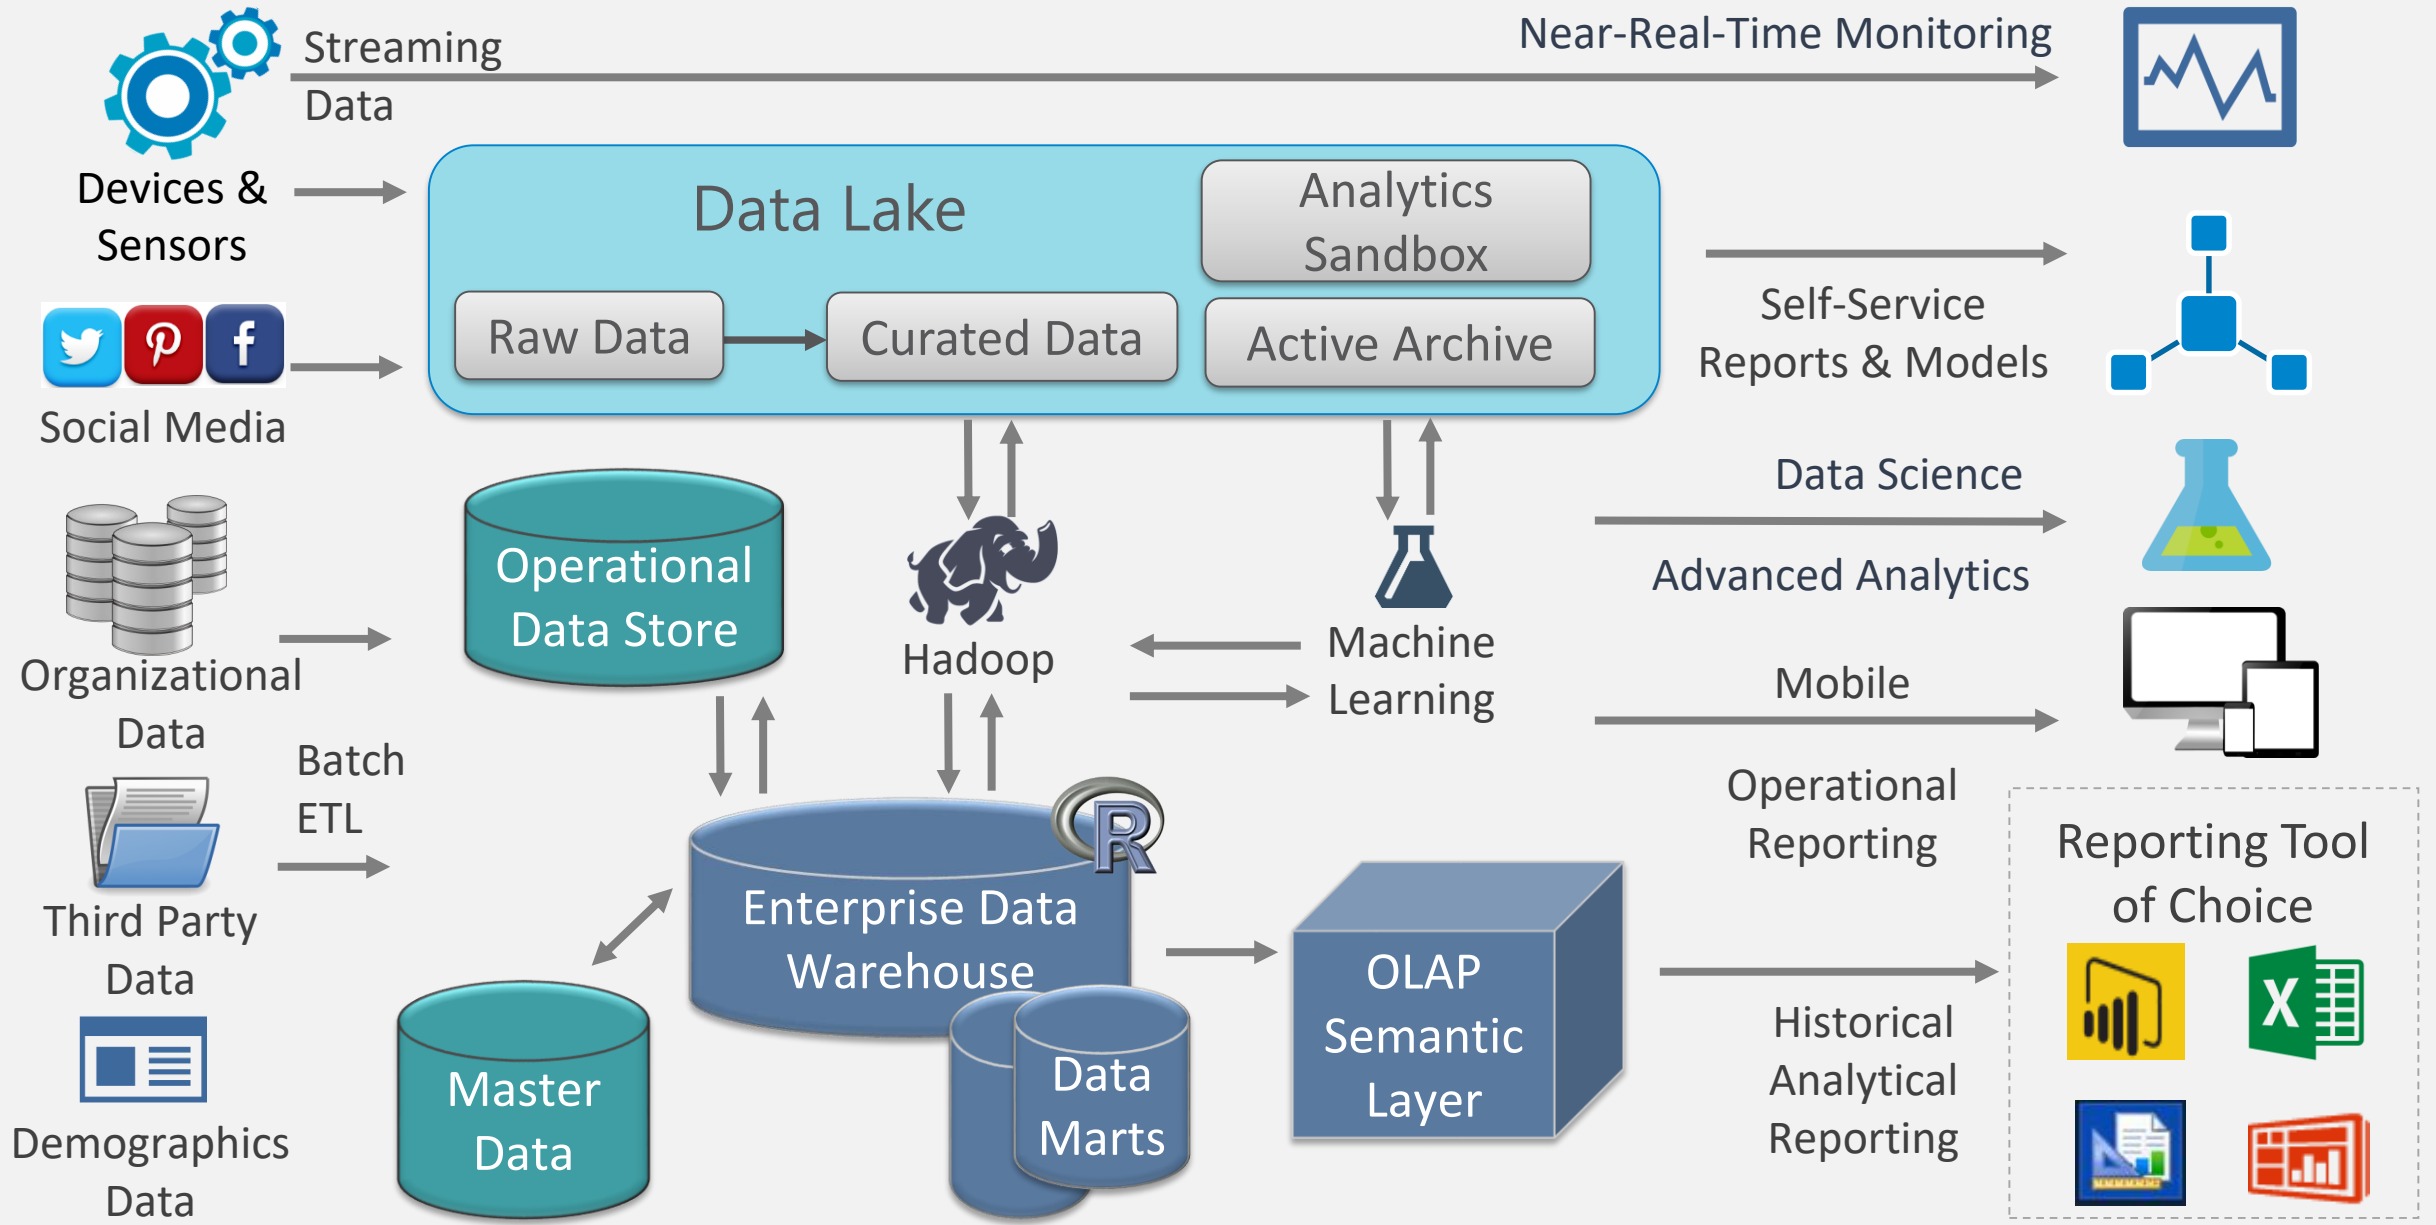
Example Objectives:

- ✓ Identify lowest-selling products
- ✓ Analyze margin per customer

DW+BI Systems Used to Be Fairly Straightforward



DW+BI Systems Have Grown in Complexity







Data Warehouse Design Principles

3 Primary Architectural Areas

Data Acquisition

Data Storage




Reporting Tool
of Choice



Data Delivery

Integrate Data from Multiple Sources



Source Systems

Objective:

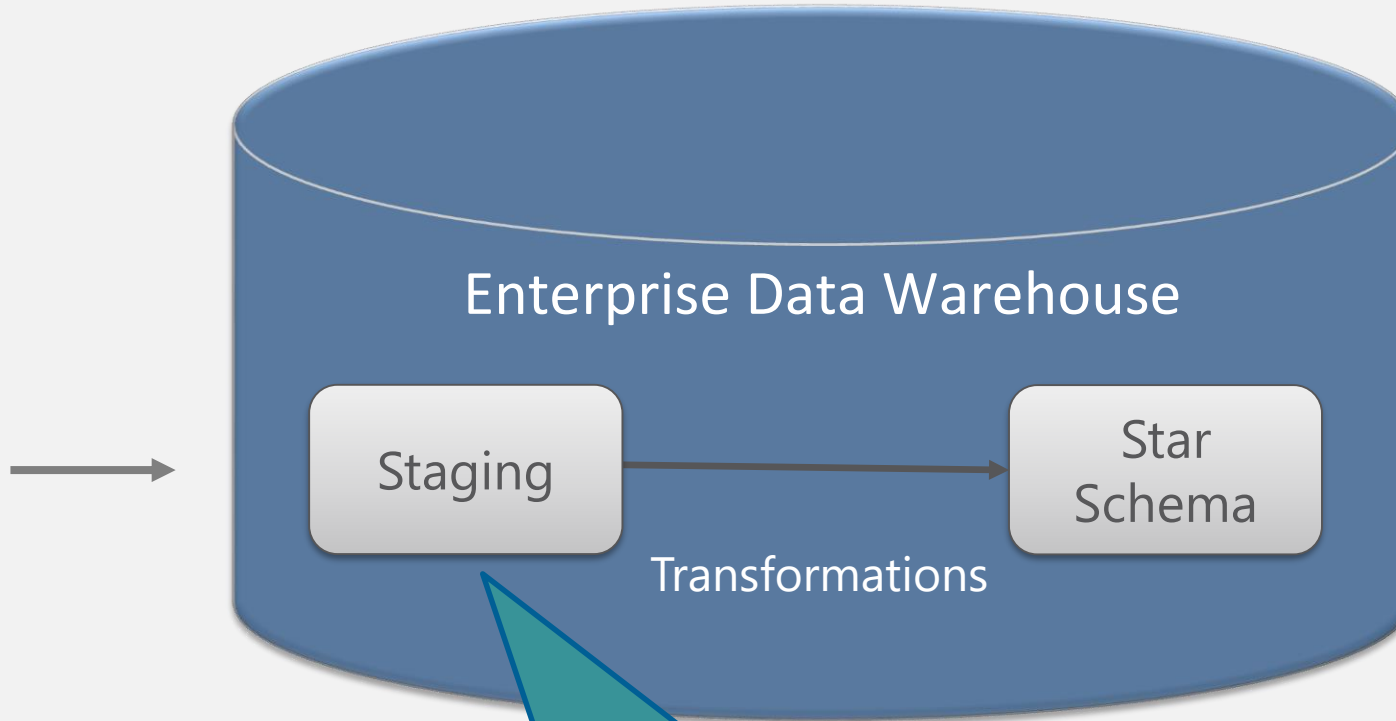

Data is inherently more **valuable** once it is integrated.

Example:

Full view of a customer:

- Sales activity +
- Delinquent invoices +
- Support/help requests

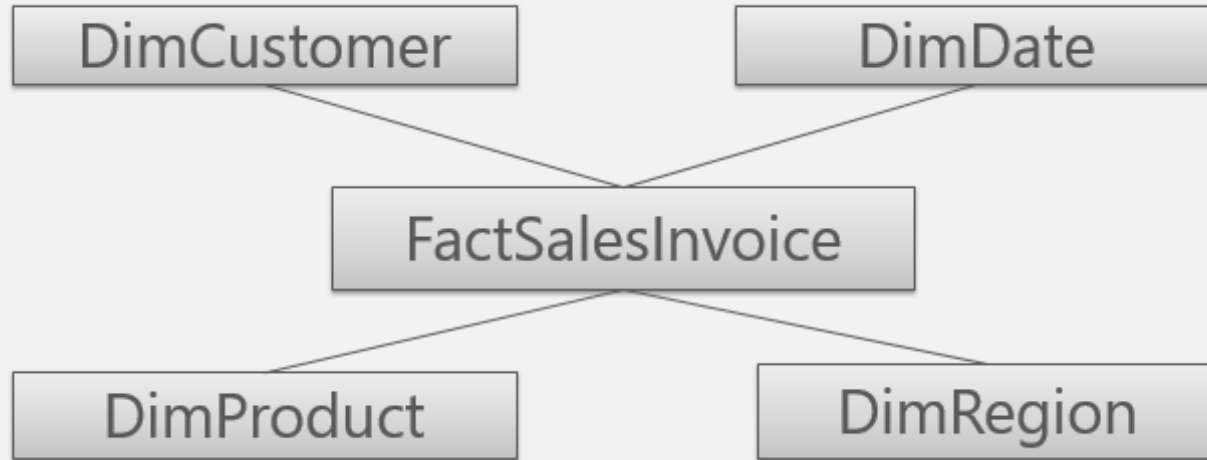
Use of Staging Environment



New trend: use of a data lake as the DW staging environment

- Staging Objectives:
- ✓ **Reduce load** on source system
 - ✓ No changes to source format
 - ✓ A "kitchen area"
 - ✓ Snapshot of source data for troubleshooting

Usage of a Star Schema



Dimension Table

Provides the **descriptive context** – attributes with the who, what, when, why, or how

Fact Table

Fact tables contain the **numeric, quantitative** data (aka **measures**)

Benefits of a Star Schema


Optimal for **known** reporting scenarios

Denormalized structure, structured around **business logic**, is good for **performance** & **consistency**

Decoupled from source systems: **surrogate keys** which have no intrinsic meaning

Usability:

- ✓ Stable, predictable environment
- ✓ Less joins, easier navigation
- ✓ Friendly, recognizable names
- ✓ History retention
- ✓ Integrate multiple systems



Challenges of a Star Schema

Requires **up-front analysis**
("schema on write")


Difficult to handle new & **unpredictable**
or exploratory scenarios

Increasing **volumes of data**

Reducing windows of time for **data loads** (near real-time is challenging)

Data quality issues are often surfaced in the reporting layer

Not practical to contain ***all* of the data** all the time



Store the Lowest Level Detail You Have

Drill-down behavior:



You may be forced to only store aggregated data for extremely high data volumes. Or, you may choose an alternative technology (like a data lake, a NoSQL database, or Hadoop).

Dimension Design

Dimension Tables

Dimension tables: provide the **descriptive context** – attributes with the who, what, when, why, or how. They should always include **friendly names & descriptions**.

Dimension tables can contain:

Type of Column in a Dim	Example
Attributes	Customer Name
Non-additive numeric value	Customer Value to Acquisition Cost Ratio
Numeric value used <i>*only*</i> for filtering or grouping (usually accompanied by a "band of ranges")	Customer Satisfaction % Customer Satisfaction Range 90%-100% 80-89% Less than 80%

*Dimension tables should ***not*** contain aggregatable numeric values (measures).*

Types of Dimension Tables

Most common types of dimensions:

Type of Dim Table	Description
Type 0	Values cannot change (ex: DimDate).
Type 1	Any value which changes is overwritten; no history is preserved.
Type 2 aka Slowly Changing Dimension	Certain important values which change generate a new row which is effective-dated. <i>(Not all columns should be type 2 - certain columns can be type 1.)</i>
Type 6	Hybrid of type 1 and 2 which includes a new column for the important values, as well as a new row.

Types 3, 4, 5, and 7 do exist, but are less commonly utilized.

Type 1 Dimension

Original
data:

Customer SK	Customer NK	Customer Name	AuditRow UpdateDate
1	ABC	Brian Jones	6-4-2014
2	DEF	Sally Baker	10-1-2015

Change to Customer Name occurs.

Updated
data:

Customer SK	Customer NK	Customer Name	AuditRow UpdateDate
1	ABC	Brian Jones	6-4-2014
2	DEF	Sally Walsh	12-2-2016

Type 2 Dimension

Original
data:

Customer SK	Customer NK	Customer Name	AuditRow Effective Date	AuditRow Expired Date	AuditRow IsCurrent
1	ABC	Brian Jones	6-4-2014	12-31-9999	1
2	DEF	Sally Baker	10-1-2015	12-31-9999	1

Change to Customer Name occurs.

Updated
data:

Customer SK	Customer NK	Customer Name	AuditRow Effective Date	AuditRow Expired Date	AuditRow IsCurrent
1	ABC	Brian Jones	6-4-2014	12-31-9999	1
2	DEF	Sally Baker	10-1-2015	12-2-2016	0
3	DEF	Sally Walsh	12-3-2016	12-31-9999	1

Type 6 Dimension

Original data:

Customer SK	Customer NK	Customer Name	Customer Name Current	AuditRow Effective Date	AuditRow Expired Date	Audit Rows Current
1	ABC	Brian Jones	Brian Jones	6-4-2014	12-31-9999	1
2	DEF	Sally Baker	Sally Baker	10-1-2015	12-31-9999	1

Change to Customer Name occurs.

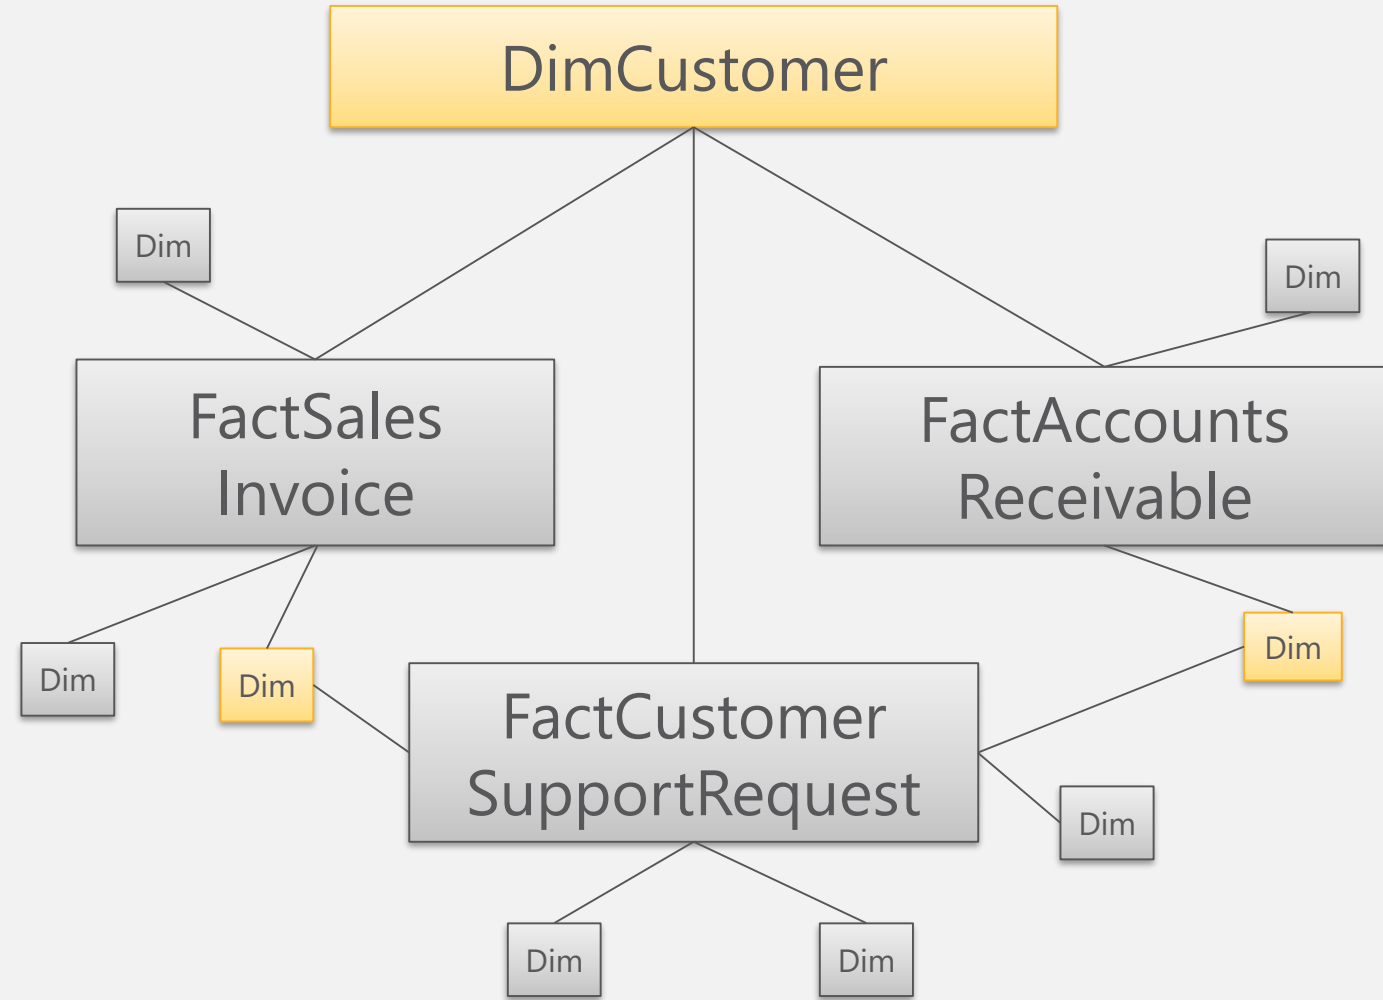
Updated data:

Customer SK	Customer NK	Customer Name	Customer Name Current	Audit Row Effective Date	AuditRow Expired Date	Audit Rows Current
1	ABC	Brian Jones	Brian Jones	6-4-2014	12-31-9999	1
2	DEF	Sally Baker	Sally Walsh	10-1-2015	12-2-2016	0
3	DEF	Sally Walsh	Sally Walsh	12-3-2016	12-31-9999	1

Conformed Dimension

A conformed dimension **reuses the same dimension** across numerous fact tables: critical for unifying data from various sources.

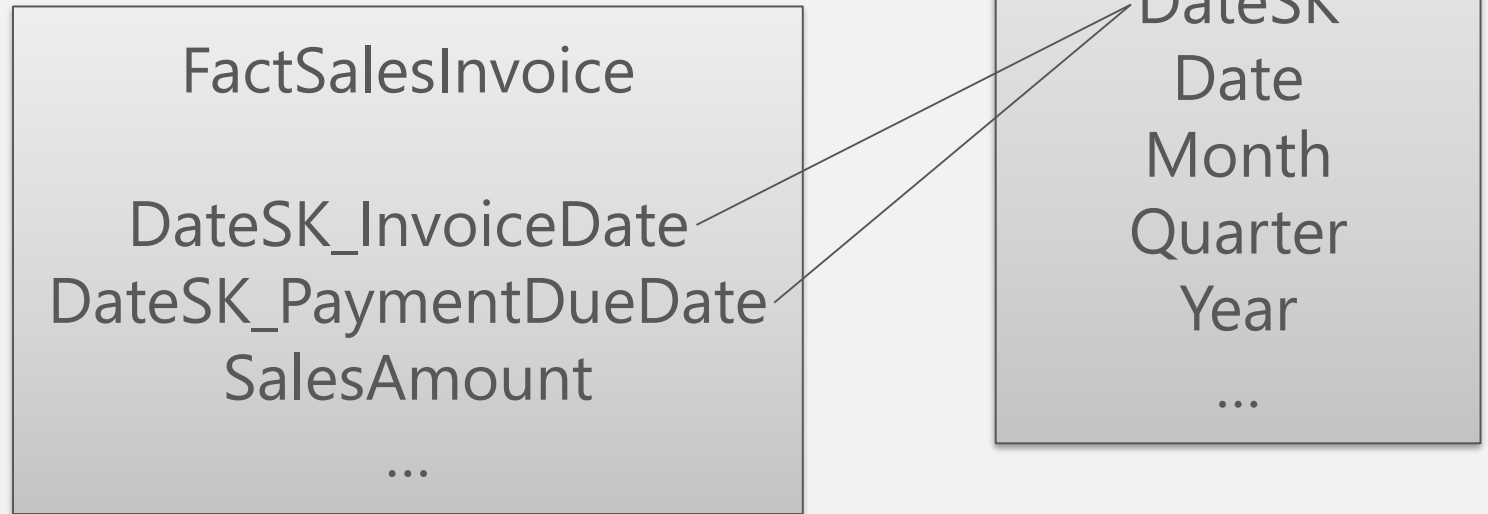
Conformed dimensions provide significant value with '**drill across**' functionality, and provide a **consistent** user experience.



Role-Playing Dimension

A role-playing dimension utilizes the same conformed dimension. Objective is to avoid creating multiple physical copies of the same dimension table.

```
SELECT
    FSI.SalesAmount
    ,InvoiceDate = DtInv.Date
    ,PymtDueDate = DtDue.Date
FROM FactSalesInvoice AS FSI
INNER JOIN DimDate AS DtInv
    ON FSI.DateSK_InvoiceDate = DtInv.DateSK
INNER JOIN DimDate AS DtDue
    ON FSI.DateSK_PaymentDueDate = DtDue.DateSK
```



Hierarchies

Hierarchies are extremely useful for handling rollups, and for drill-down & drill-through behavior.

Date Hierarchy

Year
 Quarter
 Month
 Day

Geography Hierarchy

Country
 State or Province
 City
 Address

Dimension Design

```
1 CREATE TABLE [DW].[DimCustomer] (  
2     [CustomerSK] INT IDENTITY (1, 1) NOT NULL  
3     , [RegionNumberNK] NVARCHAR(10) CONSTRAINT [dfDimCustomer_RegionNumberNK] DEFAULT (N'') NOT NULL  
4     , [CustomerNumberNK] NVARCHAR(10) CONSTRAINT [dfDimCustomer_CustomerNumberNK] DEFAULT (N'') NOT NULL  
5     , [CustomerNumber] NVARCHAR(10) CONSTRAINT [dfDimCustomer_CustomerNumber] DEFAULT (N'') NOT NULL  
6     , [CustomerName] NVARCHAR(30) CONSTRAINT [dfDimCustomer_CustomerName] DEFAULT (N'') NOT NULL  
7     , [CustomerNameCurrent] NVARCHAR(30) CONSTRAINT [dfDimCustomer_CustomerNameCurrent] DEFAULT (N'') NOT NULL  
8     , [CustomerNumberName] NVARCHAR(45) CONSTRAINT [dfDimCustomer_CustomerNumberName] DEFAULT (N'') NOT NULL  
9     , [CustomerNameNumber] NVARCHAR(45) CONSTRAINT [dfDimCustomer_CustomerNameNumber] DEFAULT (N'') NOT NULL  
10    , [CustomerFIPSCode] NVARCHAR(10) CONSTRAINT [dfDimCustomer_AccountLocationFIPSCode] DEFAULT (N'') NOT NULL  
11    , [CustomerTypeCode] NVARCHAR(10) CONSTRAINT [dfDimCustomer_CustomerTypeCode] DEFAULT (N'') NOT NULL  
12    , [CustomerTypeDescription] NVARCHAR(100) CONSTRAINT [dfDimCustomer_CustomerTypeDescription] DEFAULT (N'') NOT NULL  
13    , [CustomerTypeCode] NVARCHAR(10) CONSTRAINT [dfDimCustomer_CustomerTypeCode] DEFAULT (N'') NOT NULL  
14    , [CustomerTypeDescription] NVARCHAR(100) CONSTRAINT [dfDimCustomer_CustomerTypeDescription] DEFAULT (N'') NOT NULL
```

Inline syntax format works in the SSDT database project which requires "declarative development."

No alters beneath the create.

Dimension Design

Remove the Dim or Fact prefix from user access layers.

```
1 CREATE TABLE [DW].[DimCustomer] (  
2   [CustomerSK] INT IDENTITY (1, 1) NOT NULL  
3   , [RegionNumberNK] NVARCHAR(10) CONSTRAINT [dfDimCustomer_RegionNumberNK] DEFAULT (N'') NOT NULL  
4   , [CustomerNumberNK] NVARCHAR(10) CONSTRAINT [dfDimCustomer_CustomerNumberNK] DEFAULT (N'') NOT NULL  
5   , [CustomerNumber] NVARCHAR(10) CONSTRAINT [dfDimCustomer_CustomerNumber] DEFAULT (N'') NOT NULL  
6   , [CustomerName] NVARCHAR(30) CONSTRAINT [dfDimCustomer_CustomerName] DEFAULT (N'') NOT NULL  
7   , [CustomerNameCurrent] NVARCHAR(30) CONSTRAINT [dfDimCustomer_CustomerNameCurrent] DEFAULT (N'') NOT NULL  
8   , [CustomerNumberName] NVARCHAR(45) CONSTRAINT [dfDimCustomer_CustomerNumberName] DEFAULT (N'') NOT NULL  
9   , [CustomerNameNumber] NVARCHAR(45) CONSTRAINT [dfDimCustomer_CustomerNameNumber] DEFAULT (N'') NOT NULL  
10  , [CustomerFIPSCode] NVARCHAR(10) CONSTRAINT [dfDimCustomer_AccountLocationFIPSCode] DEFAULT (N'') NOT NULL  
11  , [CustomerTypeCode] NVARCHAR(10) CONSTRAINT [dfDimCustomer_CustomerTypeCode] DEFAULT (N'') NOT NULL  
12  , [CustomerTypeDesc] NVARCHAR(30) CONSTRAINT [dfDimCustomer_CustomerTypeDesc] DEFAULT (N'') NOT NULL  
13  , [CustomerTypeCodeDesc] NVARCHAR(36) CONSTRAINT [dfDimCustomer_CustomerTypeCodeDesc] DEFAULT (N'') NOT NULL  
14  , [CustomerTypeCodeDescCode] NVARCHAR(36) CONSTRAINT [dfDimCustomer_CustomerTypeDescCode] DEFAULT (N'') NOT NULL
```

Golden rule: a column exists in one and only one place in the DW.

Dimension Design

```
1 CREATE TABLE [DW].[DimCustomer] (  
2     [CustomerSK] INT IDENTITY (1, 1) NOT NULL  
3     , [RegionNumberNK] NVARCHAR(10) CONSTRAINT [dfDimCustomer_RegionNumberNK] FOREIGN KEY ([RegionNumberNK]) REFERENCES [RegionNumberNK] ([RegionNumberNK])  
4     , [CustomerNumberNK] NVARCHAR(10) CONSTRAINT [dfDimCustomer_CustomerNumberNK] FOREIGN KEY ([CustomerNumberNK]) REFERENCES [CustomerNumberNK] ([CustomerNumberNK])  
5     , [CustomerNumber] NVARCHAR(10) CONSTRAINT [dfDimCustomer_CustomerNumber] FOREIGN KEY ([CustomerNumber]) REFERENCES [CustomerNumber] ([CustomerNumber])  
6     , [CustomerName] NVARCHAR(30) CONSTRAINT [dfDimCustomer_CustomerName] FOREIGN KEY ([CustomerName]) REFERENCES [CustomerName] ([CustomerName])  
7     , [CustomerNameCurrent] NVARCHAR(30) CONSTRAINT [dfDimCustomer_CustomerNameCurrent] FOREIGN KEY ([CustomerNameCurrent]) REFERENCES [CustomerNameCurrent] ([CustomerNameCurrent])  
8     , [CustomerNumberName] NVARCHAR(45) CONSTRAINT [dfDimCustomer_CustomerNumberName] FOREIGN KEY ([CustomerNumberName]) REFERENCES [CustomerNumberName] ([CustomerNumberName])  
9     , [CustomerNameNumber] NVARCHAR(45) CONSTRAINT [dfDimCustomer_CustomerNameNumber] FOREIGN KEY ([CustomerNameNumber]) REFERENCES [CustomerNameNumber] ([CustomerNameNumber])  
10    , [CustomerFIPSCode] NVARCHAR(10) CONSTRAINT [dfDimCustomer_AccountLocationFIPSCode] FOREIGN KEY ([CustomerFIPSCode]) REFERENCES [CustomerFIPSCode] ([CustomerFIPSCode])  
11    , [CustomerTypeCode] NVARCHAR(10) CONSTRAINT [dfDimCustomer_CustomerTypeCode] FOREIGN KEY ([CustomerTypeCode]) REFERENCES [CustomerTypeCode] ([CustomerTypeCode])  
12    , [CustomerTypeDesc] NVARCHAR(30) CONSTRAINT [dfDimCustomer_CustomerTypeDesc] FOREIGN KEY ([CustomerTypeDesc]) REFERENCES [CustomerTypeDesc] ([CustomerTypeDesc])  
13    , [CustomerTypeCodeDesc] NVARCHAR(36) CONSTRAINT [dfDimCustomer_CustomerTypeCodeDesc] FOREIGN KEY ([CustomerTypeCodeDesc]) REFERENCES [CustomerTypeCodeDesc] ([CustomerTypeCodeDesc])  
14    , [CustomerTypeDescCode] NVARCHAR(36) CONSTRAINT [dfDimCustomer_CustomerTypeDescCode] FOREIGN KEY ([CustomerTypeDescCode]) REFERENCES [CustomerTypeDescCode] ([CustomerTypeDescCode])
```

Use a naming convention to easily identify surrogate keys & natural keys

Use the smallest datatypes you can use without risk of overflows

Make careful decisions on the use of varchar vs. nvarchar

Dimension Design

```
1 CREATE TABLE [DW].[DimCustomer] (  
2   [CustomerSK] INT IDENTITY (1, 1) NOT NULL  
3   , [RegionNumberNK] NVARCHAR(10) CONSTRAINT  
4   , [CustomerNumberNK] NVARCHAR(10) CONSTRAINT [dfDimCustomer_CustomerNumb  
5   , [CustomerNumber] NVARCHAR(10) CONSTRAINT [dfDimCustomer_CustomerNumber  
6   , [CustomerName] NVARCHAR(30) CONSTRAINT [dfDimCustomer_CustomerName] DE  
7   , [CustomerNameCurrent] NVARCHAR(30) CONSTRAINT [dfDimCustomer_CustomerN  
8   , [CustomerNumberName] NVARCHAR(45) CONSTRAINT [dfDimCustomer_CustomerNu  
9   , [CustomerNameNumber] NVARCHAR(45) CONSTRAINT [dfDimCustomer_CustomerNa  
10  , [CustomerCode] NVARCHAR(10) CONSTRAINT [dfDimCustomer_AccountLocat  
11  , [CustomerType] NVARCHAR(10) CONSTRAINT [dfDimCustomer_CustomerType  
12  , [CustomerTypeDesc] NVARCHAR(30) CONSTRAINT [dfDimCustomer_CustomerTypeDe  
13  , [CustomerTypeCodeDesc] NVARCHAR(36) CONSTRAINT [dfDimCustomer_CustomerTypeCodeDesc] DEFAULT (N'') NOT NULL  
14  , [CustomerTypeDescCode] NVARCHAR(36) CONSTRAINT [dfDimCustomer_CustomerTypeDescCode] DEFAULT (N'') NOT NULL
```

Alternatively,
could be
converted in a
view or semantic
layer. Objective is
to avoid reporting
tools trying to
sum.

Avoid numeric data types for non-
aggregatable columns such as
Customer Number.

Also useful for retaining leading 0s
or for international zip codes.

Dimension Design

```
1 CREATE TABLE [DW].[DimCustomer] (  
2   [CustomerSK] INT IDENTITY (1, 1) NOT NULL  
3   , [RegionNumberNK] NVARCHAR(10) CONSTRAINT [dfDimCustomer_RegionNumberNK] DEFAULT (N'') NOT NULL  
4   , [CustomerNumberNK] NVARCHAR(10) CONSTRAINT [dfDimCustomer_CustomerNumberNK] DEFAULT (N'') NOT NULL  
5   , [CustomerNumber] NVARCHAR(10) CONSTRAINT [dfDimCustomer_CustomerNumber] DEFAULT (N'') NOT NULL  
6   , [CustomerName] NVARCHAR(30) CONSTRAINT [dfDimCustomer_CustomerName] DEFAULT (N'') NOT NULL  
7   , [CustomerNameCurrent] NVARCHAR(30) CONSTRAINT [dfDimCustomer_CustomerNameCurrent] DEFAULT (N'') NOT NULL  
8   , [CustomerNumberName] NVARCHAR(40) CONSTRAINT [dfDimCustomer_CustomerNumberName] DEFAULT (N'') NOT NULL  
9   , [CustomerNameNumber] NVARCHAR(40) CONSTRAINT [dfDimCustomer_CustomerNameNumber] DEFAULT (N'') NOT NULL  
10  , [AccountLocationFIPSCode] NVARCHAR(10) CONSTRAINT [dfDimCustomer_AccountLocationFIPSCode] DEFAULT (N'') NOT NULL  
11  , [CustomerTypeCode] NVARCHAR(10) CONSTRAINT [dfDimCustomer_CustomerTypeCode] DEFAULT (N'') NOT NULL  
12  )
```

Default constraints are present for non-nullable columns. In a DW, defaults are optional if ETL strictly controls all data management. *Don't let SQL Server auto-name constraints.

Avoid 'Or Is Null' issues for attributes which are commonly used in predicates.

Dimension Design

```
1 CREATE TABLE [DW].[DimCustomer] (  
2   [CustomerSK] INT IDENTITY (1, 1) NOT NULL  
3   , [RegionNumberNK] NVARCHAR(10) CONSTRAINT [dfDimCustomer_RegionNumberNK] FOREIGN KEY ([RegionNumberNK]) REFERENCES [DW].[DimRegion] ([RegionSK])  
4   , [CustomerNumberNK] NVARCHAR(10) CONSTRAINT [dfDimCustomer_CustomerNumberNK] FOREIGN KEY ([CustomerNumberNK]) REFERENCES [DW].[DimCustomer] ([CustomerSK])  
5   , [CustomerNumber] NVARCHAR(10) NOT NULL  
6   , [CustomerName] NVARCHAR(30) CONSTRAINT [dfDimCustomer_CustomerName] FOREIGN KEY ([CustomerName]) REFERENCES [DW].[DimCustomer] ([CustomerSK])  
7   , [CustomerNameCurrent] NVARCHAR(30) CONSTRAINT [dfDimCustomer_CustomerNameCurrent] FOREIGN KEY ([CustomerNameCurrent]) REFERENCES [DW].[DimCustomer] ([CustomerSK]) DEFAULT (N'') NOT NULL  
8   , [CustomerNumberName] NVARCHAR(45) CONSTRAINT [dfDimCustomer_CustomerNumberName] FOREIGN KEY ([CustomerNumberName]) REFERENCES [DW].[DimCustomer] ([CustomerSK]) DEFAULT (N'') NOT NULL  
9   , [CustomerNameNumber] NVARCHAR(45) CONSTRAINT [dfDimCustomer_CustomerNameNumber] FOREIGN KEY ([CustomerNameNumber]) REFERENCES [DW].[DimCustomer] ([CustomerSK]) DEFAULT (N'') NOT NULL  
10  , [CustomerFIPSCode] NVARCHAR(10) CONSTRAINT [dfDimCustomer_AccountLocationFIPSCode] FOREIGN KEY ([CustomerFIPSCode]) REFERENCES [DW].[DimAccountLocation] ([AccountLocationSK]) DEFAULT (N'') NOT NULL  
11  , [CustomerTypeCode] NVARCHAR(10) CONSTRAINT [dfDimCustomer_CustomerTypeCode] FOREIGN KEY ([CustomerTypeCode]) REFERENCES [DW].[DimCustomerType] ([CustomerTypeSK]) DEFAULT (N'') NOT NULL  
12  , [CustomerTypeDesc] NVARCHAR(30) CONSTRAINT [dfDimCustomer_CustomerTypeDesc] FOREIGN KEY ([CustomerTypeDesc]) REFERENCES [DW].[DimCustomerType] ([CustomerTypeSK]) DEFAULT (N'') NOT NULL  
13  , [CustomerTypeCodeDesc] NVARCHAR(36) CONSTRAINT [dfDimCustomer_CustomerTypeCodeDesc] FOREIGN KEY ([CustomerTypeCodeDesc]) REFERENCES [DW].[DimCustomerType] ([CustomerTypeSK]) DEFAULT (N'') NOT NULL  
14  , [CustomerTypeDescCode] NVARCHAR(36) CONSTRAINT [dfDimCustomer_CustomerTypeDescCode] FOREIGN KEY ([CustomerTypeDescCode]) REFERENCES [DW].[DimCustomerType] ([CustomerTypeSK]) DEFAULT (N'') NOT NULL
```

When designing a Type 2 (or 6) dimension, only choose the most important columns to generate a new row when it changes.

A 'Current' column (which is the same across all rows in a Type 6 dimension) is helpful for columns commonly used in reporting so all history shows the newest value.

Dimension Design

```
1 CREATE TABLE [DW].[DimCustomer] (  
2   [CustomerSK] INT IDENTITY (1, 1) NOT NULL  
3   , [RegionNumberNK] NVARCHAR(10) CONSTRAINT [dfDimCustomer_RegionNumberNK] FOREIGN KEY ([RegionNumberNK]) REFERENCES [RegionNumberNK] ([RegionNumberNK])  
4   , [CustomerNumberNK] NVARCHAR(10) CONSTRAINT [dfDimCustomer_CustomerNumberNK] FOREIGN KEY ([CustomerNumberNK]) REFERENCES [CustomerNumberNK] ([CustomerNumberNK])  
5   , [CustomerNumber] NVARCHAR(10) CONSTRAINT [dfDimCustomer_CustomerNumber] FOREIGN KEY ([CustomerNumber]) REFERENCES [CustomerNumber] ([CustomerNumber])  
6   , [CustomerName] NVARCHAR(30) CONSTRAINT [dfDimCustomer_CustomerName] FOREIGN KEY ([CustomerName]) REFERENCES [CustomerName] ([CustomerName])  
7   , [CustomerNameCurrent] NVARCHAR(30) CONSTRAINT [dfDimCustomer_CustomerNameCurrent] FOREIGN KEY ([CustomerNameCurrent]) REFERENCES [CustomerNameCurrent] ([CustomerNameCurrent])  
8   , [CustomerNumberName] NVARCHAR(45) CONSTRAINT [dfDimCustomer_CustomerNumberName] FOREIGN KEY ([CustomerNumberName]) REFERENCES [CustomerNumberName] ([CustomerNumberName])  
9   , [CustomerNameNumber] NVARCHAR(45) CONSTRAINT [dfDimCustomer_CustomerNameNumber] FOREIGN KEY ([CustomerNameNumber]) REFERENCES [CustomerNameNumber] ([CustomerNameNumber])  
10  , [CustomerFIPSCode] NVARCHAR(10) CONSTRAINT [dfDimCustomer_AccountLocationFIPSCode] FOREIGN KEY ([CustomerFIPSCode]) REFERENCES [AccountLocationFIPSCode] ([AccountLocationFIPSCode])  
11  , [CustomerTypeCode] NVARCHAR(10) CONSTRAINT [dfDimCustomer_CustomerTypeCode] FOREIGN KEY ([CustomerTypeCode]) REFERENCES [CustomerTypeCode] ([CustomerTypeCode])  
12  , [CustomerTypeDesc] NVARCHAR(100) CONSTRAINT [dfDimCustomer_CustomerTypeDesc] FOREIGN KEY ([CustomerTypeDesc]) REFERENCES [CustomerTypeDesc] ([CustomerTypeDesc])  
13  , [CustomerTypeCodeDesc] NVARCHAR(100) CONSTRAINT [dfDimCustomer_CustomerTypeCodeDesc] FOREIGN KEY ([CustomerTypeCodeDesc]) REFERENCES [CustomerTypeCodeDesc] ([CustomerTypeCodeDesc])  
14  , [CustomerTypeDescCode] NVARCHAR(100) CONSTRAINT [dfDimCustomer_CustomerTypeDescCode] FOREIGN KEY ([CustomerTypeDescCode]) REFERENCES [CustomerTypeDescCode] ([CustomerTypeDescCode])  
15  )
```

Could also be derived in views or semantic layer. Or, computed columns could be used.

Optionally, can store variations of concatenated columns such as:

Name (Number)

Number - Name

Description (Code)

Code - Description

Dimension Design

```
45 , [CustomerDeliveryOnSaturday] NVARCHAR(50) NOT NULL
46 , [AuditETLBatchID] INT NOT NULL
47 , [AuditInsertDate] DATETIME CONSTRAINT [dfDimCustomer_AuditInsertDate] FOREIGN KEY ([AuditETLBatchID]) REFERENCES [AuditETLBatchID]
48 , [AuditInsertBy] NVARCHAR(50) NOT NULL CONSTRAINT [dfDimCustomer_AuditInsertBy] FOREIGN KEY ([AuditInsertBy]) REFERENCES [AuditInsertBy]
49 , [AuditHashValue] BINARY(20) CONSTRAINT [dfDimCustomer_AuditHashValue] FOREIGN KEY ([AuditHashValue]) REFERENCES [AuditHashValue]
50 , [AuditModifiedDate] DATETIME NULL
51 , [AuditModifiedBy] NVARCHAR(50) NULL
52 , [AuditIsDeleted] BIT CONSTRAINT [dfDimCustomer_AuditIsDeleted] FOREIGN KEY ([AuditIsDeleted]) REFERENCES [AuditIsDeleted]
53 , [AuditIsPurgeEligible] BIT CONSTRAINT [dfDimCustomer_AuditIsPurgeEligible] FOREIGN KEY ([AuditIsPurgeEligible]) REFERENCES [AuditIsPurgeEligible]
54 , [AuditRowEffectiveDate] DATETIME CONSTRAINT [dfDimCustomer_RowEffectiveDate] FOREIGN KEY ([AuditRowEffectiveDate]) REFERENCES [AuditRowEffectiveDate]
55 , [AuditRowExpiredDate] DATETIME CONSTRAINT [dfDimDimCustomer_RowExpiredDate] FOREIGN KEY ([AuditRowExpiredDate]) REFERENCES [AuditRowExpiredDate] DEFAULT ('12/31/2999') NOT NULL
56 , [AuditRowIsCurrent] BIT CONSTRAINT [dfDimDimCustomer_RowIsCurrent] FOREIGN KEY ([AuditRowIsCurrent]) REFERENCES [AuditRowIsCurrent] DEFAULT ((0)) NOT NULL
57 , CONSTRAINT [pkCustomerSK] PRIMARY KEY CLUSTERED ([CustomerSK] ASC) ON [Dimensions]
58 , CONSTRAINT [uqDimCustomer] UNIQUE NONCLUSTERED ([CustomerNumberNK] ASC,
59 , [RegionNumberNK] ASC,
60 , [AuditRowEffectiveDate] ASC,
61 , [AuditRowExpiredDate] ASC) ON [Dimensions]
62 ) ON [Dimensions]
```

Standard audit columns.
The 'Audit' prefix makes it clear they are generated in the DW not the source.

Additional columns if the Type 2 historical change tracking is occurring.

Dimension Design

```
45 , [CustomerDeliveryOnSaturday] NVARCHAR(3) CONSTRAINT [dfDimCustomer_DeliveryOnSaturday]
46 , [AuditETLBatchID] INT NOT NULL CONSTRAINT [dfDimCustomer_AuditETLBatchID]
47 , [AuditInsertDate] DATETIME CONSTRAINT [dfDimCustomer_AuditInsertDate]
48 , [AuditInsertBy] NVARCHAR(50) NOT NULL CONSTRAINT [dfDimCustomer_AuditInsertBy]
49 , [AuditHashValue] BINARY(20) CONSTRAINT [dfDimCustomer_AuditHashValue]
50 , [AuditModifiedDate] DATETIME NULL
51 , [AuditModifiedBy] NVARCHAR(50) NULL
52 , [AuditIsDeleted] BIT CONSTRAINT [dfDimCustomer_AuditIsDeleted] DEFAULT ((0)) NOT NULL
53 , [AuditIsPurgeEligible] BIT CONSTRAINT [dfDimCustomer_AuditIsPurgeEligible] DEFAULT ((0)) NOT NULL
54 , [AuditRowEffectiveDate] DATETIME CONSTRAINT [dfDimCustomer_RowEffectiveDate] DEFAULT (SYSDATETIME()) NOT NULL
55 , [AuditRowExpiredDate] DATETIME CONSTRAINT [dfDimCustomer_RowExpiredDate] DEFAULT ('12/31/2999') NOT NULL
56 , [AuditRowIsCurrent] BIT CONSTRAINT [dfDimCustomer_RowIsCurrent] DEFAULT ((0)) NOT NULL
57 , CONSTRAINT [pkCustomerSK] PRIMARY KEY CLUSTERED ([CustomerSK] ASC) ON [Dimensions]
58 , CONSTRAINT [uqDimCustomer] UNIQUE NONCLUSTERED ([CustomerNumberNK] ASC,
59 , [RegionNumberNK] ASC,
60 , [AuditRowEffectiveDate] ASC,
61 , [AuditRowExpiredDate] ASC) ON [Dimensions]
62 ) ON [Dimensions];
```

All key & index suggestions are merely a starting point. As your DW grows, you might have to refine your strategy depending on ETL.

Primary key based on the surrogate key. This is also our clustered index.

Dimension Design

```
45 , [CustomerDeliveryOnSaturday] NVARCHAR(3) CONSTRAINT [dfD  
46 , [AuditETLBatchID] INT NOT NULL CONSTRAINT [dfDimCustomer  
47 , [AuditInsertDate] DATETIME CONSTRAINT [dfDimCustomer_Aud  
48 , [AuditInsertBy] NVARCHAR(50) NOT NULL CONSTRAINT [dfDimC  
49 , [AuditHashValue] BINARY(20) CONSTRAINT [dfDimCustomer_Aud  
50 , [AuditModifiedDate] DATETIME NULL  
51 , [AuditModifiedBy] NVARCHAR(50) NULL  
52 , [AuditIsDeleted] BIT CONSTRAINT [dfDimCustomer_AuditIsDelete ((0)) NOT NULL  
53 , [AuditIsPurgeEligible] BIT CONSTRAINT [dfDimCustomer_AuditIsPurgeEligible] DEFAULT ((0)) NOT NULL  
54 , [AuditRowEffectiveDate] DATETIME CONSTRAINT [dfDimCustomer_AuditRowEffectiveDate] DEFAULT (SYSDATETIME()) NOT NULL  
55 , [AuditRowExpiredDate] DATETIME CONSTRAINT [dfDimCustomer_AuditRowExpiredDate] DEFAULT ('12/31/2999') NOT NULL  
56 , [AuditRowIsCurrent] BIT CONSTRAINT [dfDimCustomer_AuditRowIsCurrent] DEFAULT ((0)) NOT NULL  
57 , CONSTRAINT [pkCustomerSK] PRIMARY KEY CLUSTERED ([CustomerSK] ASC) ON [Dimensions]  
58 , CONSTRAINT [uqDimCustomer] UNIQUE NONCLUSTERED ([CustomerNumberNK] ASC,  
59 [RegionNumberNK] ASC,  
60 [AuditRowEffectiveDate] ASC,  
61 [AuditRowExpiredDate] ASC) ON [Dimensions]  
62 ) ON [Dimensions];
```

The unique constraint implicitly creates a unique index as well, which will assist with lookup operations.

Unique constraint, based on natural keys, defines the “grain” of the table. It also helps identify data quality issues & is very helpful to the SQL Server query optimizer.

Dimension Design

Use of non-Primary filegroups.
Ex: Dimensions, Facts,
Staging, Other.

```
45 , [CustomerDeliveryOnSaturday] NVARCHAR(3) CONSTRAINT  
46 , [AuditETLBatchID] INT NOT NULL CONSTRAINT [dfDimCust  
47 , [AuditInsertDate] DATETIME CONSTRAINT [dfDimCustomer  
48 , [AuditInsertBy] NVARCHAR(50) NOT NULL CONSTRAINT [df  
49 , [AuditHashValue] BINARY(20) CONSTRAINT [dfDimCustomer_AuditHash  
50 , [AuditModifiedDate] DATETIME NULL  
51 , [AuditModifiedBy] NVARCHAR(50) NULL  
52 , [AuditIsDeleted] BIT CONSTRAINT [dfDimCustomer_AuditIsDeleted] DEFAULT NOT NULL  
53 , [AuditIsPurgeEligible] BIT CONSTRAINT [dfDimCustomer_AuditIsPurgeEligible] DEFAULT ((0)) NOT NULL  
54 , [AuditRowEffectiveDate] DATETIME CONSTRAINT [dfDimCustomer_RowEffectiveDate] DEFAULT (SYSDATETIME()) NOT NULL  
55 , [AuditRowExpiredDate] DATETIME CONSTRAINT [dfDimDimCustomer_RowExpiredDate] DEFAULT ('12/31/2999') NOT NULL  
56 , [AuditRowIsCurrent] BIT CONSTRAINT [dfDimDimCustomer_RowIsCurrent] DEFAULT ((0)) NOT NULL  
57 , CONSTRAINT [pkCustomerSK] PRIMARY KEY CLUSTERED ([CustomerSK] ASC) ON [Dimensions]  
58 , CONSTRAINT [uqDimCustomer] UNIQUE NONCLUSTERED ([CustomerNumberNK] ASC,  
59 [RegionNumberNK] ASC,  
60 [AuditRowEffectiveDate] ASC,  
61 [AuditRowExpiredDate] ASC) ON [Dimensions]  
62 ) ON [Dimensions];
```

Fact Design

Fact Tables

Fact tables contain the **numeric, quantitative** data (aka **measures**).

Typically **one fact table per distinct business process**.

Exception: "consolidated" facts (aka "merged" facts) such as actual vs. forecast which require the same granularity and are frequently analyzed together.

Fact tables can contain:

Type of Column in a Fact	Example
Measures	Sales Amount
Foreign keys to dimension table	3392 (meaningless integer surrogate key)
Degenerate dimension	Order Number

Types of Fact Tables

Most common types of facts:

Type of Fact Table	Description	Example
Transaction Fact	An event at a point in time	FactSalesInvoice
Periodic Snapshot Fact	Summary at a point in time	FactARBalanceDaily
Accumulating Snapshot Fact	Summary across the lifetime of an event	FactStudentApplication
Timespan Tracking Fact	Effective-dated rows	FactCapitalAssetBalance

Other facts:

Type of Fact Table	Description	Example
Factless Fact Table	Recording when an event did not occur	FactPromotionNoSales
Aggregate Facts	Rollups, usually to improve reporting speed	FactSalesInvoiceSummary

Fact Design

```
1 CREATE TABLE DW.FactSalesInvoice(  
2     SalesInvoiceSK INT  
3     CONSTRAINT dfFactSalesInvoice_ARObligation  
4     CONSTRAINT fkFactSalesInvoice_DimSalesInvoice REFERENCES DW.DimSalesInvoice(SalesInvoiceSK)  
5     ,CustomerSK INT  
6     CONSTRAINT dfFactSalesInvoice_CustomerSK DEFAULT ((-1)) NOT NULL  
7     CONSTRAINT fkFactSalesInvoice_DimCustomer REFERENCES DW.DimCustomer(CustomerSK)  
8     ,RegionSK SMALLINT  
9     CONSTRAINT dfFactSalesInvoice_RegionSK DEFAULT ((-1)) NOT NULL  
10    CONSTRAINT fkFactSalesInvoice_DimRegion REFERENCES DW.DimRegion(RegionSK)  
11    ,DateSK_AROpenedDate INT  
12    CONSTRAINT dfFactSalesInvoice_DateSK_AROpenedDate DEFAULT ((-1)) NOT NULL  
13    CONSTRAINT fkFactSalesInvoice_DimDate_AROpenedDate REFERENCES DW.DimDate(DateSK)  
14    ,DateSK_ARClosedDate INT  
15    CONSTRAINT dfFactSalesInvoice_DateSK_ARClosedDate DEFAULT(29991231) NOT NULL  
16    CONSTRAINT fkFactSalesInvoice_DimDate_ARClosedDate REFERENCES DW.DimDate(DateSK)  
17    ,DateSK_ARDiscountDate INT  
18    CONSTRAINT dfFactSalesInvoice_DateSK_ARDiscountDate DEFAULT(29991231) NOT NULL  
19    CONSTRAINT fkFactSalesInvoice_DimDate_ARDiscountDate REFERENCES DW.DimDate(DateSK)
```

Even if all of the SKs are the same, avoid combining fact tables for unrelated business processes.

One fact table per distinct business process.

Fact Design

The combination of SKs might dictate the grain of the fact table, but it may not.

```
1 CREATE TABLE DW.FactSalesInvoice
2 SalesInvoiceSK INT
3 CONSTRAINT dfFactSalesInvoice_ARObligationSK DEFAULT ((-1)) NOT NULL
4 CONSTRAINT fkFactSalesInvoice_DimSalesInvoice REFERENCES DW.DimSalesInvoice(SalesInvoiceSK)
5 ,CustomerSK INT
6 CONSTRAINT dfFactSalesInvoice_CustomerSK DEFAULT ((-1)) NOT NULL
7 CONSTRAINT fkFactSalesInvoice_DimCustomer REFERENCES DW.DimCustomer(CustomerSK)
8 ,RegionSK SMALLINT
9 CONSTRAINT dfFactSalesInvoice_RegionSK DEFAULT ((-1)) NOT NULL
10 CONSTRAINT fkFactSalesInvoice_DimRegion REFERENCES DW.DimRegion(RegionSK)
11 ,DateSK_AROpenedDate INT
12 CONSTRAINT dfFactSalesInvoice_DateSK_AROpenedDate DEFAULT(19000101) NOT NULL
13 CONSTRAINT fkFactSalesInvoice_DimDate_AROpenedDate REFERENCES DW.DimDate(DateSK)
14 ,DateSK_ARClosedDate INT
15 CONSTRAINT dfFactSalesInvoice_DateSK_ARClosedDate DEFAULT(29991231) NOT NULL
16 CONSTRAINT fkFactSalesInvoice_DimDate_ARClosedDate REFERENCES DW.DimDate(DateSK)
17 ,DateSK_ARDiscountDate INT
18 CONSTRAINT dfFactSalesInvoice_DateSK_ARDiscountDate DEFAULT(29991231) NOT NULL
19 CONSTRAINT fkFactSalesInvoice_DimDate_ARDiscountDate REFERENCES DW.DimDate(DateSK)
```

Fact Design

Some data modelers prefer the unknown member row to have its key assigned randomly.

```
1 CREATE TABLE DW.FactSalesInvoice
2 SalesInvoiceSK INT
3     CONSTRAINT dfFactSalesInvoice_ARObligationSK DEFAULT ((-1)) NOT NULL
4     CONSTRAINT fkFactSalesInvoice_DimSalesInvoice REFERENCES DW.DimSalesInvoice(SalesInvoiceSK)
5 ,CustomerSK INT
6     CONSTRAINT dfFactSalesInvoice_CustomerSK DEFAULT ((-1)) NOT NULL
7     CONSTRAINT fkFactSalesInvoice_DimCustomer REFERENCES DW.DimCustomer(CustomerSK)
8 ,RegionSK SMALLINT
9     CONSTRAINT dfFactSalesInvoice_RegionSK DEFAULT ((-1)) NOT NULL
10    CONSTRAINT fkFactSalesInvoice_DimRegion REFERENCES DW.DimRegion(RegionSK)
11 ,DateSK_AROpenedDate INT
12    CONSTRAINT dfFactSalesInvoice_DateSK_AROpenedDate DEFAULT(19000101) NOT NULL
13    CONSTRAINT fkFactSalesInvoice_DimDate_AROpenedDate REFERENCES DW.DimDate(DateSK)
14 ,DateSK_ARClosedDate INT
15    CONSTRAINT dfFactSalesInvoice_DateSK_ARClosedDate DEFAULT(29991231) NOT NULL
16    CONSTRAINT fkFactSalesInvoice_DimDate_ARClosedDate REFERENCES DW.DimDate(DateSK)
17 ,DateSK_ARDiscountDate INT
18    CONSTRAINT dfFactSalesInvoice_DateSK_ARDiscountDate DEFAULT(19000101) NOT NULL
19    CONSTRAINT fkFactSalesInvoice_DimDate_ARDiscountDate REFERENCES DW.DimDate(DateSK)
```

Default equates to the 'unknown member' row.

Fact Design

```
1 CREATE TABLE DW.FactSalesInvoice(  
2     SalesInvoiceSK INT  
3     CONSTRAINT dfFactSalesInvoice_AR  
4     CONSTRAINT fkFactSalesInvoice_Di  
5     ,CustomerSK INT  
6     CONSTRAINT dfFactSalesInvoice_Cu  
7     CONSTRAINT fkFactSalesInvoice_Dim  
8     ,RegionSK SMALLINT  
9     CONSTRAINT dfFactSalesInvoice_RegionSK DEFAULT ( ) NOT NULL  
10    CONSTRAINT fkFactSalesInvoice_DimRegion REFERENCES DW.DimRegion(RegionSK)  
11    ,DateSK_AROpenedDate INT  
12    CONSTRAINT dfFactSalesInvoice_DateSK_AROpenedDate DEFAULT(19000101) NOT NULL  
13    CONSTRAINT fkFactSalesInvoice_DimDate_AROpenedDate REFERENCES DW.DimDate(DateSK)  
14    ,DateSK_ARClosedDate INT  
15    CONSTRAINT dfFactSalesInvoice_DateSK_ARClosedDate DEFAULT(29991231) NOT NULL  
16    CONSTRAINT fkFactSalesInvoice_DimDate_ARClosedDate REFERENCES DW.DimDate(DateSK)  
17    ,Dat  
18  
19
```

Optionally can use two types of Date defaults: one in the past, one in the future. Helps with 'Less than' or 'Greater than' predicates.

It's also fine for a date SK to be an actual date datatype instead of an integer.

Fact Design

Having a PK in a fact is personal preference. Usually you don't want a clustered index on it though.

```
1 CREATE TABLE DW.FactSalesInvoice(  
2     SalesInvoiceSK INT  
3     CONSTRAINT dfFactSalesInvoice_ARObligationSK DEFAULT ((-1)) NOT NULL  
4     CONSTRAINT fkFactSalesInvoice_DimSalesInvoice REFERENCES DW.DimSalesInvoice(SalesInvoiceSK)  
5     ,CustomerSK INT  
6     CONSTRAINT dfFactSalesInvoice_CustomerSK DEFAULT ((-1)) NOT NULL  
7     CONSTRAINT fkFactSalesInvoice_DimCustomer REFERENCES DW.DimCustomer(CustomerSK)  
8     ,RegionSK SMALLINT  
9     CONSTRAINT dfFactSalesInvoice_RegionSK DEFAULT ((-1)) NOT NULL  
10    CONSTRAINT fkFactSalesInvoice_DimRegion REFERENCES DW.DimRegion(RegionSK)  
11    ,DateSK_AROpenedDate INT  
12    CONSTRAINT dfFactSalesInvoice_DateSK_AROpenedDate DEFAULT(19000101) NOT NULL  
13    CONSTRAINT fkFactSalesInvoice_DimDate_AROpenedDate REFERENCES DW.DimDate(DateSK)  
14    ,DateSK_ARClosedDate INT  
15    CONSTRAINT dfFactSalesInvoice_DateSK_ARClosedDate DEFAULT(29991231) NOT NULL  
16    CONSTRAINT fkFactSalesInvoice_DimDate_ARClosedDate REFERENCES DW.DimDate(DateSK)  
17    ,DateSK_ARStartDate INT  
18    CONSTRAINT dfFactSalesInvoice_DateSK_ARStartDate DEFAULT(29991231) NOT NULL  
19    CONSTRAINT fkFactSalesInvoice_DimDate_ARStartDate REFERENCES DW.DimDate(DateSK)
```

Foreign key constraints mitigate referential integrity issues.

Fact Design

```
22
23 ,WarehouseSK INT
24     CONSTRAINT dfFactSalesInvoice_WarehouseSK DEFAULT ((-1)) NOT NULL
25     CONSTRAINT fkFactSalesInvoice_DimWarehouse REFERENCES DW.DimWarehouse(WarehouseSK)
26 ,ProductSK INT
27     CONSTRAINT dfFactSalesInvoice_ProductSK DEFAULT ((-1)) NOT NULL
28     CONSTRAINT fkFactSalesInvoice_DimProduct REFERENCES DW.DimProduct(ProductSK)
29 ,EmployeeSK_SalespersonForTransaction INT
30     CONSTRAINT dfFactSalesInvoice_EmployeeSK_SalespersonForTransaction DEFAULT ((-1)) NOT NULL
31     CONSTRAINT fkFactSalesInvoice_DimEmployee_SalespersonForTransaction REFERENCES DW.DimEmployee(EmployeeSK)
32 ,NumberOfInvoices INT NULL
33 ,TotalInvoiceAmount DECIMAL(10,2) NULL
34 ,GrossSalesAmount DECIMAL(10,2) NULL
35 ,FreightAmount DECIMAL(10,2) NULL
36 ,TaxAmount DECIMAL(10,2) NULL
37 ,NetSalesAmount DECIMAL(10,2) NULL
38 ,DiscountAmount DECIMAL(10,2) NULL
39 ,InvoiceNumberNK NVARCHAR(10) CONSTRAINT dfFactSalesInvoice_InvoiceNumberNK NOT NULL
40 ,InvoiceItemNumberNK NVARCHAR(10) CONSTRAINT dfFactSalesInvoice_InvoiceItemNumberNK NOT NULL
    ,EmployeeSK_SalespersonForTransaction NVARCHAR(10) CONSTRAINT dfFactSalesInvoice_EmployeeSK_SalespersonForTransaction NOT NULL
```

Measures are sparse,
therefore nullable.
0s are not stored except in a
factless fact table.

Fact Design

```
47 ,WarehouseNumberNK NVARCHAR(10) CONSTRAINT dfFactSalesInvoice_WarehouseNumberNK DEFAULT (N'') NOT NULL
48 ,RegionNumberNK NVARCHAR(3) CONSTRAINT dfFactSalesInvoice_RegionNumberNK DEFAULT (N'') NOT NULL
49 ,AuditETLBatchID INT NOT NULL CONSTRAINT dfFactSalesInvoice_AuditKey DEFAULT ((0))
50 ,AuditInsertDate DATETIME
51 ,AuditInsertBy NVARCHAR(50)
52 ,AuditHashValue BINARY(16)
53 ,AuditModifiedDate DATETIME
54 ,AuditModifiedBy NVARCHAR(50)
55 ,AuditIsDeleted BIT CONSTRAINT [ucFactSalesInvoice_AuditIsDeleted] UNIQUE
56 ,AuditIsPurgeEligible BIT
57 ,CONSTRAINT [uqFactSalesInvoice] UNIQUE ([CustomerNumberNK], [InvoiceNumberNK], [EmployeeSK_SalespersonForTransaction])
58 ([CustomerNumberNK], [InvoiceNumberNK], [EmployeeSK_SalespersonForTransaction])
59 ,INDEX [ixFactSalesInvoice] ON [FactSalesInvoice] ([CustomerNumberNK], [InvoiceNumberNK], [EmployeeSK_SalespersonForTransaction])
60 WITH (FILLFACTOR = 80)
61 ,INDEX [ixFactSalesInvoice] ON [FactSalesInvoice] ([CustomerNumberNK], [InvoiceNumberNK])
62 ,INDEX [ixFactSalesInvoice] ON [FactSalesInvoice] ([InvoiceNumberNK], [EmployeeSK_SalespersonForTransaction])
63 ,INDEX [ixFactSalesInvoice] ON [FactSalesInvoice] ([EmployeeSK_SalespersonForTransaction])
64 ,INDEX [ixFactSalesInvoice_WarehouseNumberNK] ON [FactSalesInvoice] ([WarehouseNumberNK])
65 ,INDEX [ixFactSalesInvoice_EmployeeSK_SalespersonForTransaction] NONCLUSTERED ([EmployeeSK_SalespersonForTransaction] ASC)
66 WITH (FILLFACTOR = 80)
67 ,INDEX [ixFactSalesInvoice_CustomerNumberNK] ON [FactSalesInvoice] ([CustomerNumberNK])
68 ,INDEX [ixFactSalesInvoice_InvoiceNumberNK] ON [FactSalesInvoice] ([InvoiceNumberNK])
69 ,INDEX [ixFactSalesInvoice_EmployeeSK_SalespersonForTransaction] ON [FactSalesInvoice] ([EmployeeSK_SalespersonForTransaction])
70 ) ON [FactSalesInvoice]
```

Natural key in a fact violates Kimball rules.
However, they are helpful for:

- (1) Re-assigning SK if a lookup issue occurred and an unknown member got assigned.
- (2) Allows unique constraint on the NKs for ensuring data integrity.

****Never (ever!) let the NKs be exposed or used for anything besides ETL. And only create minimum # of NKs to identify the row.****

```
80) ON [FactSalesInvoice]
ON [FactSalesInvoice]
ON [FactSalesInvoice]
```

Fact Design

```
47 ,WarehouseNumberNK NVARCHAR(10) CONSTRAINT dfFactSalesInvoice
48 ,RegionNumberNK NVARCHAR(3) CONSTRAINT dfFactSalesInvoice_F
49 ,AuditETLBatchID INT NOT NULL CONSTRAINT dfFactSalesInvoice
50 ,AuditInsertDate DATETIME CONSTRAINT dfFactSalesInvoice_Au
51 ,AuditInsertBy NVARCHAR(50) NOT NULL CONSTRAINT dfFactSales
52 ,AuditHashValue BINARY(20) CONSTRAINT dfFactSalesInvoice_Au
53 ,AuditModifiedDate DATETIME NULL
54 ,AuditModifiedBy NVARCHAR(50) NULL
55 ,AuditIsDeleted BIT CONSTRAINT dfFactSalesInvoice_AuditIsDel DEFAULT ((0)) NOT NULL
56 ,AuditIsPurgeEligible BIT CONSTRAINT dfFactSalesInvoice_AuditIsPurgeEligible DEFAULT ((0)) NOT NULL
57 ,CONSTRAINT [uqFactSalesInvoice] UNIQUE NONCLUSTERED
58 ([CustomerNumberNK] ASC, [InvoiceNumberNK] ASC, [InvoiceItemNumberNK] ASC) ON [Facts]
59 ,INDEX [ixFactSalesInvoice_DateSK_AROpenedDate] CLUSTERED (DateSK_AROpenedDate ASC)
60 WITH (FILLFACTOR = 80, COMPRESSION = PAGE) ON [Facts]
61 ,INDEX [ixFactSalesInvoice_InvoiceSK] NONCLUSTERED (SalesInvoiceSK ASC) WITH (FILLFACTOR = 80) ON [Facts]
62 ,INDEX [ixFactSalesInvoice_CustomerSK] NONCLUSTERED (CustomerSK ASC) WITH (FILLFACTOR = 80) ON [Facts]
63 ,INDEX [ixFactSalesInvoice_RegionSK] NONCLUSTERED (RegionSK ASC) WITH (FILLFACTOR = 80) ON [Facts]
64 ,INDEX [ixFactSalesInvoice_WarehouseSK] NONCLUSTERED (WarehouseSK ASC) WITH (FILLFACTOR = 80) ON [Facts]
65 ,INDEX [ixFactSalesInvoice_TransactionSK] NONCLUSTERED (TransactionSK ASC)
66 WITH (FILLFACTOR = 80) ON [Facts]
67 ,INDEX [ixFactSalesInvoice_TransactionSK] NONCLUSTERED (TransactionSK ASC) WITH (FILLFACTOR = 80) ON [Facts]
68 ,INDEX [ixFactSalesInvoice_TransactionSK] NONCLUSTERED (TransactionSK ASC) WITH (FILLFACTOR = 80) ON [Facts]
69 ,INDEX [ixFactSalesInvoice_TransactionSK] NONCLUSTERED (TransactionSK ASC) WITH (FILLFACTOR = 80) ON [Facts]
70 ) ON [Facts] ;
```

The unique constraint implicitly creates a unique index as well, which will assist with lookup operations.

Unique constraint, based on natural keys, defines the "grain" of the table & helps identify data quality issues.

Fact Design

The clustered index is usually on a date.

```
47 , WarehouseNumberNK NVARCHAR(10) CONSTRAINT dfFactSalesInvoice_WarehouseNumberNK NOT NULL
48 , RegionNumberNK NVARCHAR(3) CONSTRAINT dfFactSalesInvoice_RegionNumberNK NOT NULL
49 , AuditETLBatchID INT NOT NULL CONSTRAINT dfFactSalesInvoice_AuditETLBatchID NOT NULL
50 , AuditInsertDate DATETIME CONSTRAINT dfFactSalesInvoice_AuditInsertDate NOT NULL
51 , AuditInsertBy NVARCHAR(50) NOT NULL CONSTRAINT dfFactSalesInvoice_AuditInsertBy SNAME()
52 , AuditHashValue BINARY(20) CONSTRAINT dfFactSalesInvoice_AuditHashValue DEFAULT ((0)) NOT NULL
53 , AuditModifiedDate DATETIME NULL
54 , AuditModifiedBy NVARCHAR(50) NULL
55 , AuditIsDeleted BIT CONSTRAINT dfFactSalesInvoice_AuditIsDeleted DEFAULT ((0)) NOT NULL
56 , AuditIsPurgeEligible BIT CONSTRAINT dfFactSalesInvoice_AuditIsPurgeEligible DEFAULT ((0)) NOT NULL
57 , CONSTRAINT [uqFactSalesInvoice] UNIQUE NONCLUSTERED
58     ([CustomerNumberNK] ASC, [InvoiceNumberNK] ASC, [InvoiceItemNumberNK] ASC) ON [Facts]
59 , INDEX [ixFactSalesInvoice_DateSK_AROpenedDate] CLUSTERED (DateSK_AROpenedDate ASC)
60     WITH (FILLFACTOR = 80, DATA_COMPRESSION = PAGE) ON [Facts]
61 , INDEX [ixFactSalesInvoice_SalesInvoiceSK] NONCLUSTERED (SalesInvoiceSK ASC) WITH (FILLFACTOR = 80) ON [Facts]
62 , INDEX [ixFactSalesInvoice_CustomerSK] NONCLUSTERED (CustomerSK ASC) WITH (FILLFACTOR = 80) ON [Facts]
63 , INDEX [ixFactSalesInvoice_RegionSK] NONCLUSTERED (RegionSK ASC) WITH (FILLFACTOR = 80) ON [Facts]
64 , INDEX [ixFactSalesInvoice_WarehouseSK] NONCLUSTERED (WarehouseSK ASC) WITH (FILLFACTOR = 80) ON [Facts]
65 , INDEX [ixFactSalesInvoice_EmployeeSK_SalespersonForTransaction] NONCLUSTERED (EmployeeSK_SalespersonForTransaction ASC)
66     WITH (FILLFACTOR = 80) ON [Facts]
67 , INDEX [ixFactSalesInvoice_DateSK_AROpenedDate] NONCLUSTERED (DateSK_AROpenedDate ASC) WITH (FILLFACTOR = 80) ON [Facts]
68 , INDEX [ixFactSalesInvoice_DateSK_AROpenedDate] NONCLUSTERED (DateSK_AROpenedDate ASC) WITH (FILLFACTOR = 80) ON [Facts]
69 , INDEX [ixFactSalesInvoice_DateSK_AROpenedDate] NONCLUSTERED (DateSK_AROpenedDate ASC) WITH (FILLFACTOR = 80) ON [Facts]
70 ) ON [Facts] ;
```

Compression set on the clustered index rather than the table.

Fact Design

```
47 ,WarehouseNumberNK NVARCHAR(10) CONSTRAINT dfFactSal
48 ,RegionNumberNK NVARCHAR(3) CONSTRAINT dfFactSalesIn
49 ,AuditETLBatchID INT NOT NULL CONSTRAINT dfFactSales
50 ,AuditInsertDate DATETIME CONSTRAINT dfFactSalesInv
51 ,AuditInsertBy NVARCHAR(50) NOT NULL CONSTRAINT dfFa
52 ,AuditHashValue BINARY(20) CONSTRAINT dfFactSalesInv
53 ,AuditModifiedDate DATETIME NULL
54 ,AuditModifiedBy NVARCHAR(50) NULL
55 ,AuditIsDeleted BIT CONSTRAINT dfFactSalesInvoice_Aud
56 ,AuditIsPurgeEligible BIT CONSTRAINT dfFactSalesInvoice_
57 ,CONSTRAINT [uqFactSalesInvoice] UNIQUE NONCLUSTERED
58 ([CustomerNumberNK] ASC, [InvoiceNumberNK] ASC, [Inv
59 ,INDEX [ixFactSalesInvoice_DateSK_AROpenedDate] CLUS (DateSK_AROpenedDate ASC)
60 WITH (FILLFACTOR = 80, DATA_COMPRESSION = PAGE) ON [Facts]
61 ,INDEX [ixFactSalesInvoice_SalesInvoiceSK] NONCLUSTERED (SalesInvoiceSK ASC) WITH (FILLFACTOR = 80) ON [Facts]
62 ,INDEX [ixFactSalesInvoice_CustomerSK] NONCLUSTERED (CustomerSK ASC) WITH (FILLFACTOR = 80) ON [Facts]
63 ,INDEX [ixFactSalesInvoice_RegionSK] NONCLUSTERED (RegionSK ASC) WITH (FILLFACTOR = 80) ON [Facts]
64 ,INDEX [ixFactSalesInvoice_WarehouseSK] NONCLUSTERED (WarehouseSK ASC) WITH (FILLFACTOR = 80) ON [Facts]
65 ,INDEX [ixFactSalesInvoice_EmployeeSK_SalespersonForTransaction] NONCLUSTERED (EmployeeSK_SalespersonForTransaction ASC)
66 WITH (FILLFACTOR = 80) ON [Facts]
67 ,INDEX [ixFactSalesInvoice_DateSK_ARDiscountDate] NONCLUSTERED (DateSK_ARDiscountDate ASC) WITH (FILLFACTOR = 80) ON [Facts]
68 ,INDEX [ixFactSalesInvoice_DateSK_ARNetDueDate] NONCLUSTERED (DateSK_ARNetDueDate ASC) WITH (FILLFACTOR = 80) ON [Facts]
69 ,INDEX [ixFactSalesInvoice_DateSK_ARClosedDate] NONCLUSTERED (DateSK_ARClosedDate ASC) WITH (FILLFACTOR = 80) ON [Facts]
70 ) ON [Facts] ;
```

Nonclustered index on each surrogate key. Useful for smaller fact tables (which don't justify a clustered columnstore index).

When to Use Columnstore Indexes or Partitioning

Handling Larger Fact Tables

Clustered Columnstore Index

Useful for:

- ✓ Reducing **data storage** due to compression of redundant values
- ✓ Improving **query** times for large datasets
- ✓ Improving **query** times due to reduced I/O (ex: column elimination)

Table Partitioning

Useful for:

- ✓ Improving **data load** times due to partition switching
- ✓ Flexibility for **maintenance** on larger tables
- ✓ Improving **query** performance (*possibly*) due to parallelism & partition elimination behavior

Clustered Columnstore Index

Simplified & conceptual

Rowstore:

Page

CounterName	Disk	DateMeasurementTaken	TimeMeasurementTaken	Measurement
Avg. Disk sec/Read	G:\	1/30/2017	4:48:41 PM	0.01818
Avg. Disk sec/Read	L:\	1/30/2017	4:48:41 PM	0.00385
Avg. Disk sec/Read	T:\	1/30/2017	4:48:41 PM	0.00780
Avg. Disk Bytes/Read	G:\	1/30/2017	4:48:41 PM	53120.73782
Avg. Disk Bytes/Read	L:\	1/30/2017	4:48:41 PM	42362.51095
Avg. Disk Bytes/Read	T:\	1/30/2017	4:48:41 PM	47951.40657

Columnstore:

Page

CounterName
Avg. Disk sec/Read
Avg. Disk Bytes/Read

Page

Disk
G:\
L:\
T:\

Page

DateMeasurementTaken
1/30/2017

Page

TimeMeasurementTaken
4:48:41 PM

Page

Measurement
0.01818
0.00385
0.00780
53120.73782
42362.51095
47951.40657

Reduced storage for low cardinality columns

Clustered Columnstore Index

Simplified & conceptual

CCI most suitable for:

✓ Tables **over 1 million rows**

Page

CounterName
Avg. Disk sec/Read
Avg. Disk Bytes/Read

Page

Disk
G:\
L:\
T:\

Page

DateMeasurementTaken
1/30/2017

Page

TimeMeasurementTaken
4:48:41 PM

Page

Measurement
0.01818
0.00385
0.00780
53120.73782
42362.51095
47951.40657

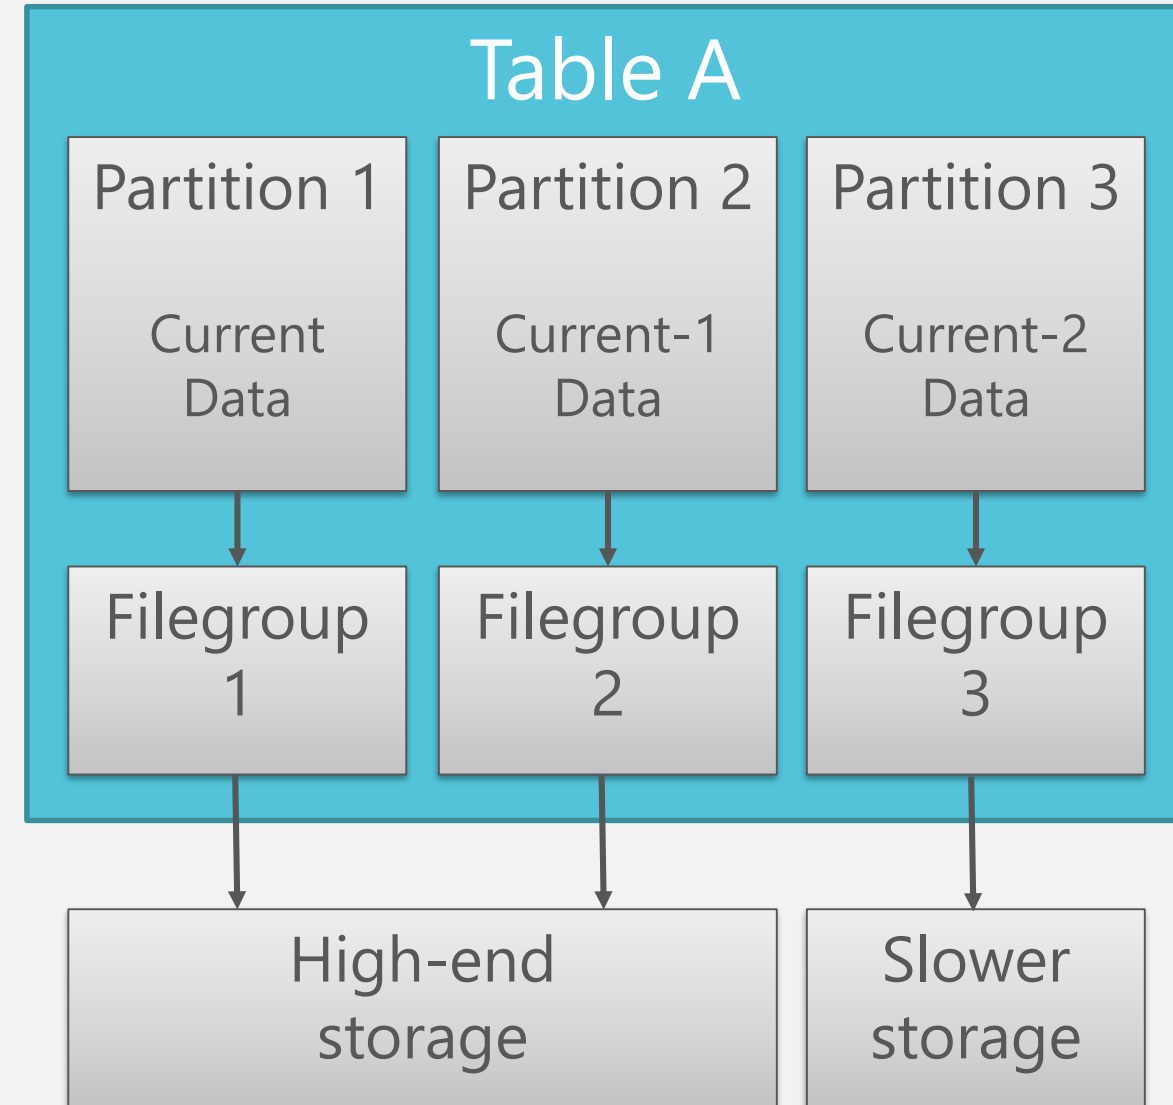
- ✓ Data structured in a **denormalized** star schema format (DW not OLTP)
- ✓ Support for **analytical query workload** which scans a large number of rows, and retrieves few columns
- ✓ Data which is **not frequently updated** ('cold' data not 'hot')
- ✓ Can selectively be used on insert-oriented workloads (ex: IoT)

(A nonclustered columnstore index targets analytical queries on an OLTP rather than a data warehouse.)

Partitioned Table

Useful for:

- ✓ **Speeding up ETL processes**
 - ✓ Large datasets (50GB+)
 - ✓ Small maintenance windows
 - ✓ Use of a sliding window
- ✓ **Storage of partitions on separate drives (filegroups)**
 - ✓ Older (cold) data on cheaper storage
 - ✓ Historical data on read-only filegroup
- ✓ **Speeding up queries (possibly)**
 - ✓ Partition elimination
 - ✓ Parallelism



Partitioned View

Useful for:

- ✓ Query performance (similar to partitioned table)
- ✓ Sharing of a single table ("partition") across multiple views
- ✓ Displaying info from > 1 database or server (via a linked server)

Requires "Check" constraints on the underlying tables (usually on a date column)


```
1 CREATE VIEW DW.vwFactSales
2 WITH SCHEMABINDING
3 AS
4
5 SELECT SalesInvoiceSK
6        ,CustomerSK
7        ,DateSK_AROpenedDate
8        ,TotalInvoiceAmount
9        ,GrossSalesAmount
10 FROM DW.FactSalesCurrent
11
12 UNION ALL
13
14 SELECT SalesInvoiceSK
15        ,CustomerSK
16        ,DateSK_AROpenedDate
17        ,TotalInvoiceAmount
18        ,GrossSalesAmount
19 FROM DW.FactSalesHistory;
20
```


Data Warehouse Tips

Handling Many-to-Many Scenarios

Classic many-to-many scenarios:

- ✓ A sales order is for many products, and a product is on many sales orders
- ✓ A customer has multiple bank accounts, and a bank account belongs to multiple customers



Ways to Track History in a DW

Most common options for tracking history:

1. Slowly changing dimension
2. Fact snapshot tables
3. Timestamp tracking fact

New option in SQL Server 2016:

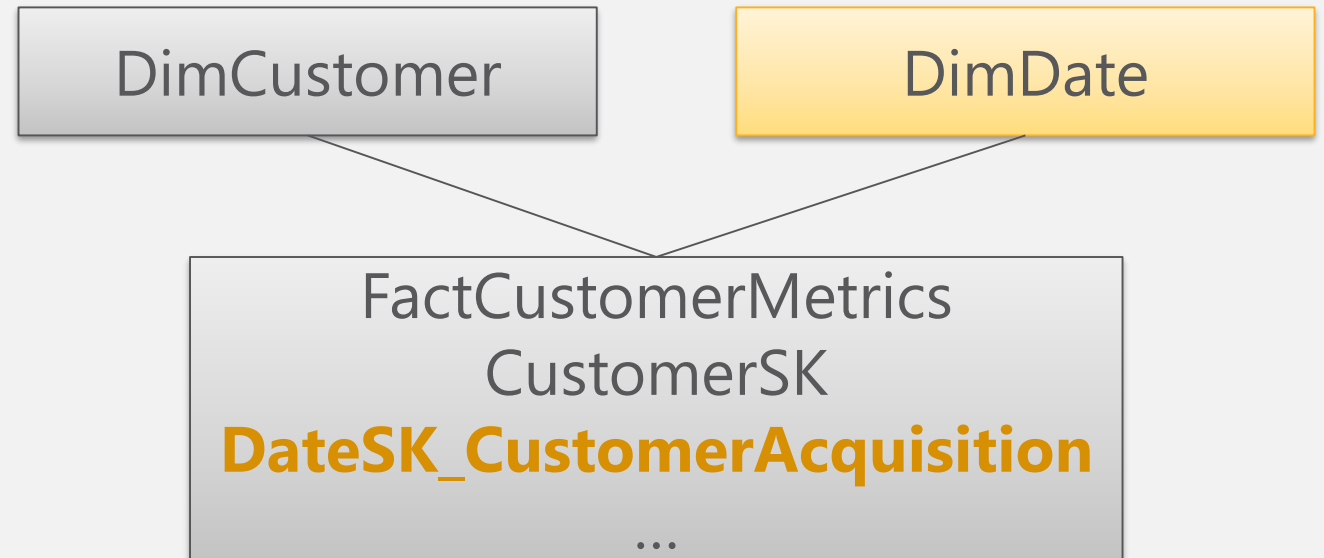
4. Temporal data tables → *Not a full replacement for slowly changing dimensions, but definitely useful for auditing*

"Smart Dates" vs. "Dumb Dates" in a DW

A "dumb date" is just an attribute:



A "smart date" relates to a full-fledged Date dimension which allows significant time analysis capabilities:



Handling of Nulls in Dimensions

Rule of thumb is to **avoid nulls** in attribute columns.

Remember the
NOT NULL and
default
constraints

What happens with this:

```
SELECT CustomerType WHERE CustomerType <> 'Retail'
```

Too easy to forget:

```
SELECT CustomerType WHERE CustomerType <> 'Retail'  
OR CustomerType IS NULL
```

Handling of Nulls in Facts

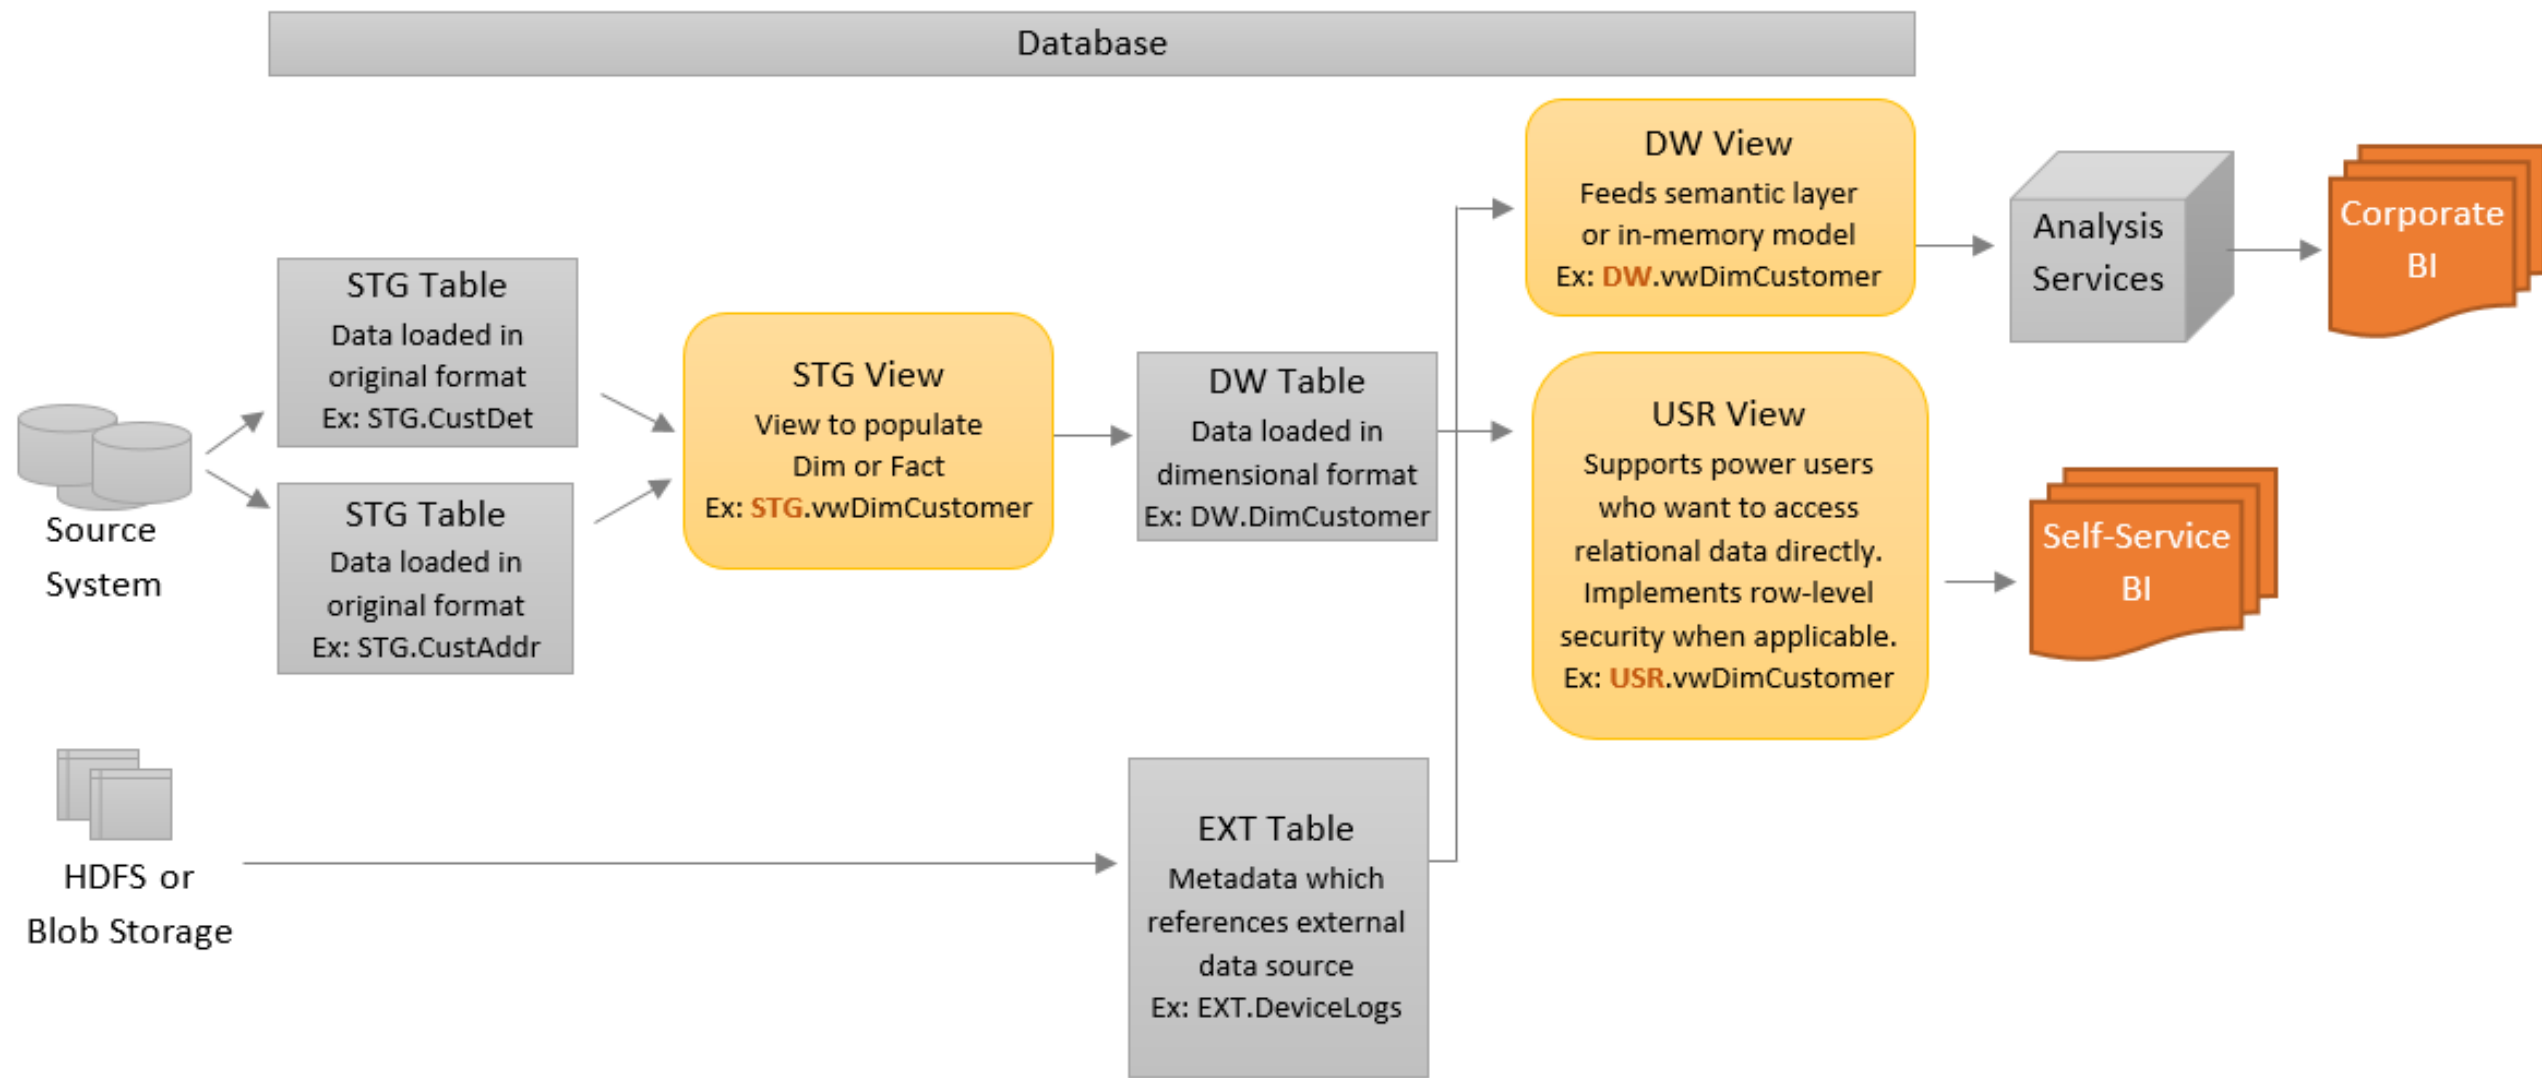
Best practice is to **avoid nulls in foreign keys**. (However, nulls are ok for a measure.)

By using an 'unknown member' relationship to the dimension, you can:

- ✓ Safely do inner joins
- ✓ Allow the fact record to be inserted & meet referential integrity
- ✓ Allow the fact record to be inserted which avoids understating measurement amounts

Ex: Just because one key is unknown, such as an EmployeeSK for who rang up the sale, should the sale not be counted?

Views Customized for Different Purposes



Recap of Important DW Design Principles

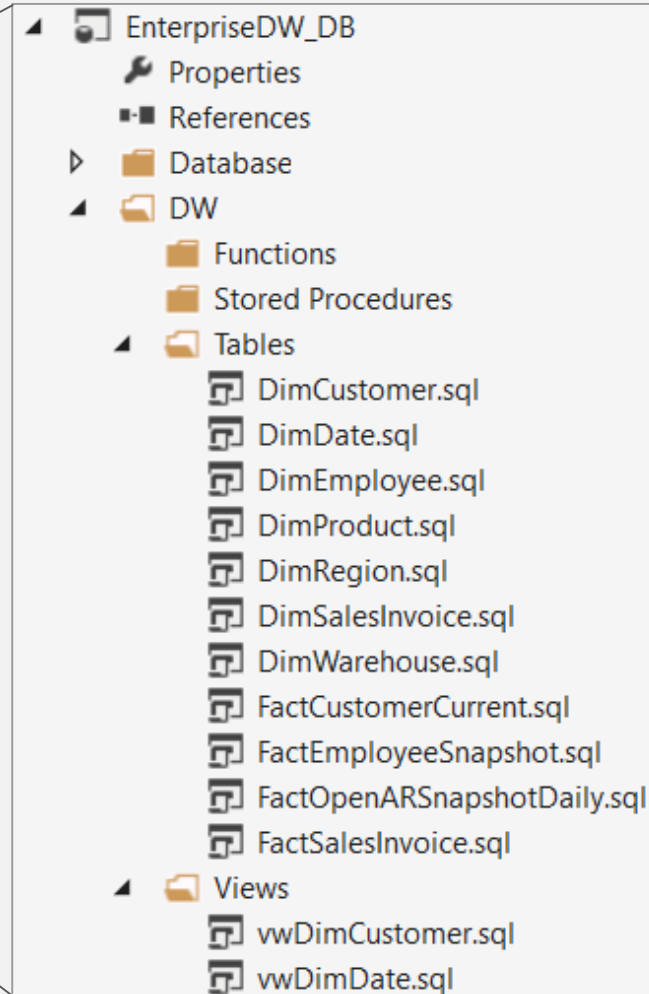

- ✓ Staging as a “kitchen” area
- ✓ Integrate data from multiple systems to increase its value
- ✓ Denormalize the data into a star schema
- ✓ A column exists in one and only one place in the star schema
- ✓ Avoid snowflake design most of the time
- ✓ Use surrogate keys which are independent from source systems
- ✓ Use conformed dimensions
- ✓ Know the grain of every table
- ✓ Have a strategy for handling changes, and for storage of history
- ✓ Store the lowest level of detail that you can
- ✓ Use an ‘unknown member’ to avoid understating facts
- ✓ Transform the data, but don’t “fix” it in the DW
- ✓ Structure your dimensional model around business processes

Recap of Important DW Design Principles

- ✓ Design facts around a single business event
- ✓ Always use friendly names & descriptions
- ✓ Use an explicit date dimension in a "role-playing" way
- ✓ Utilize bridge tables to handle many-to-many scenarios
- ✓ Plan for complexities such as:
 - ✓ Header/line data
 - ✓ Semi-additive facts
 - ✓ Multiple currencies
 - ✓ Multiple units of measure
 - ✓ Alternate hierarchies and calculations per business units
 - ✓ Allocation of measures in a snowflake design
 - ✓ Reporting of what didn't occur (factless facts)
 - ✓ Dimensional only analysis

SSDT “Database Project” Tips

Database Project Format



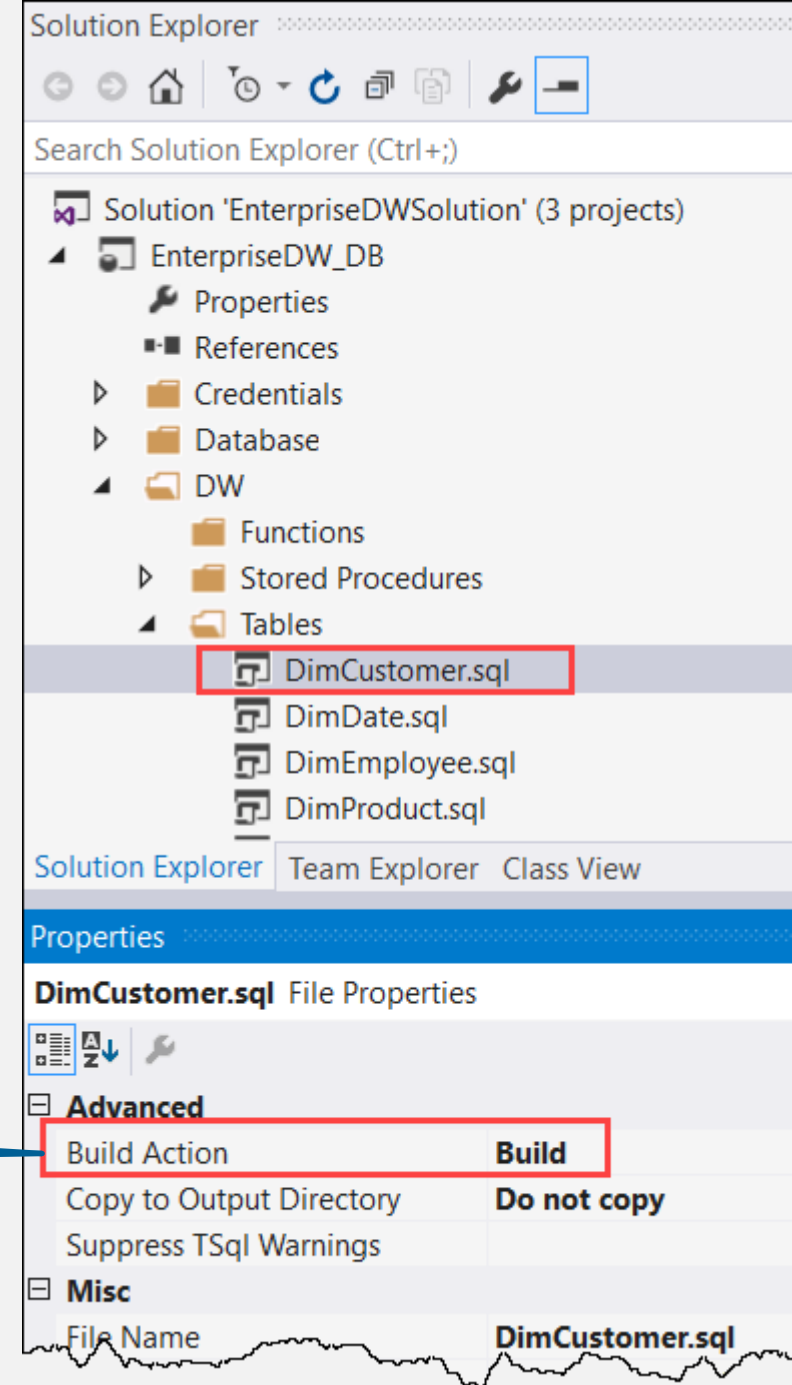
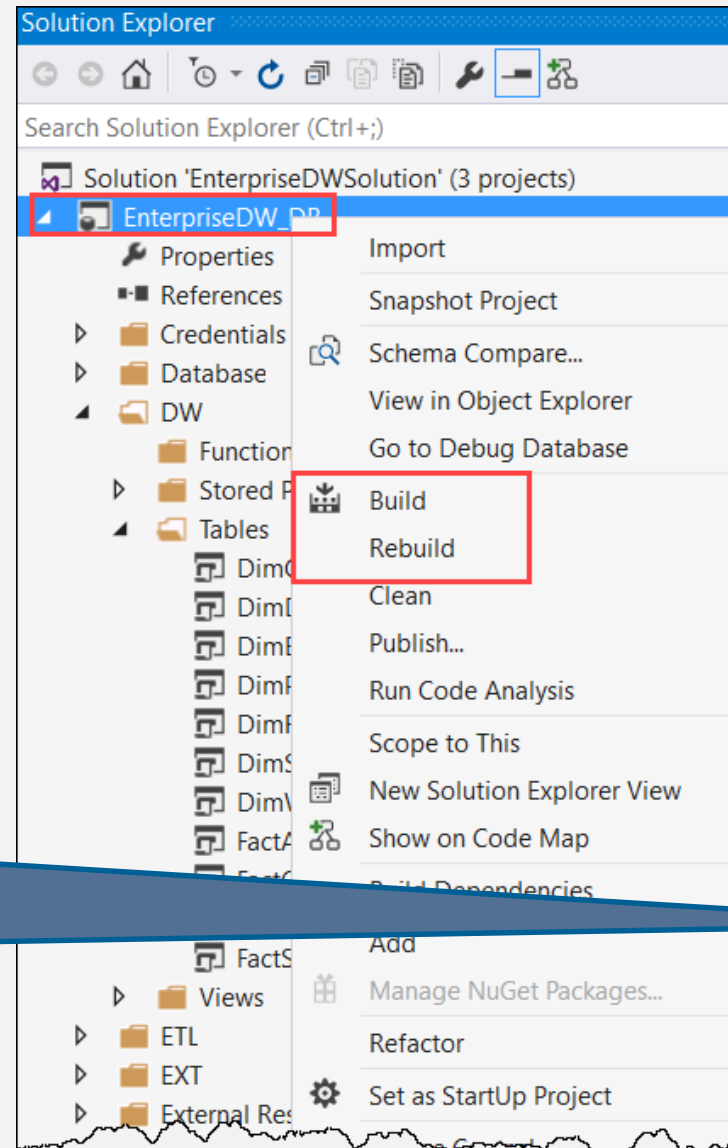
This project is organized by:

- 1 – Schema
(or Category)
- 2 – Object Type
- 3 – Object

Building the Database Project

Build frequently
to verify no
errors or
missing
references

Nearly all
objects should
be set to Build



Database Design

Pre-sized files

Auto-grow
allowed in sizeable
increments
(just in case)

Separate disks
to locate data
& log

```
1  USE [master]
2  GO
3
4  CREATE DATABASE [EnterpriseDW]
5  CONTAINMENT = NONE
6  ON PRIMARY
7  ( NAME = N'EnterpriseDW', FILENAME = N'G:\MSSQL\Data\EnterpriseDW',
8    SIZE = 64MB , MAXSIZE = UNLIMITED, FILEGROWTH = 256MB ),
9
10 FILEGROUP [Dimensions] DEFAULT
11 ( NAME = N'EnterpriseDW_dimensions', FILENAME = N'G:\MSSQL\Data\EnterpriseDW_dimensions',
12   SIZE = 3GB , MAXSIZE = UNLIMITED, FILEGROWTH = 256MB ),
13
14 FILEGROUP [Facts]
15 ( NAME = N'EnterpriseDW_facts', FILENAME = N'G:\MSSQL\Data\EnterpriseDW_facts.ndf' ,
16   SIZE = 3GB , MAXSIZE = UNLIMITED, FILEGROWTH = 256MB ),
17
18 FILEGROUP [Staging]
19 ( NAME = N'EnterpriseDW_staging', FILENAME = N'G:\MSSQL\Data\EnterpriseDW_staging.ndf' ,
20   SIZE = 2GB , MAXSIZE = UNLIMITED, FILEGROWTH = 256MB ),
21
22 FILEGROUP [Other]
23 ( NAME = N'EnterpriseDW_other', FILENAME = N'G:\MSSQL\Data\EnterpriseDW_other.ndf' ,
24   SIZE = 1GB , MAXSIZE = UNLIMITED, FILEGROWTH = 256MB )
25
26 LOG ON
27 ( NAME = N'EnterpriseDW_log', FILENAME = N'L:\MSSQL\Log\EnterpriseDW_log.ldf' ,
28   SIZE = 256MB , MAXSIZE = 2048GB , FILEGROWTH = 256MB )
```

Unknown Member Row

The SK reference in a fact table if the real value is unknown or does not exist.

Build action =
none since this is
DML

--Step 1. Permit an explicit value to be inserted into identity column

```
SET IDENTITY_INSERT [DW].[DimCustomer] ON;  
GO
```

--Step 2. Insert unknown member row

```
INSERT INTO [DW].[DimCustomer]  
    ([CustomerSK]  
    , [RegionNumberNK]  
    , [CustomerNumberNK]  
    , [CustomerNumber]  
    , [CustomerName]  
    , [AuditRowIsCurrent])  
VALUES  
    (-1  
    , N'Unknown'  
    , N'Unknown'  
    , N'Unknown'  
    , N'Unknown'  
    , '1901-01-01'  
    , '2999-12-31'  
    , 1  
    );  
GO
```

Identity_Insert does
require elevated
permissions

--Step 3. Disable ability for an explicit value to be inserted into identity column

```
SET IDENTITY_INSERT [DW].[DimCustomer] OFF;  
GO
```

Manually Maintained Data

Maintain a DML script in a Lookup (LKP) table instead of hard-coding in the ETL.






```
INSERT [LKP].[SalesInvoiceOrderType] (  
[OrderTypeCode], [OrderTypeDescription], [OrderTypeChannel], [AuditUpdateDate], [AuditUpdateBy])  
VALUES ('C ', 'Warehouse Credit', 'Warehouse', GETDATE(), SUSER_SNAME() )  
GO  
INSERT [LKP].[SalesInvoiceOrderType] (  
[OrderTypeCode], [OrderTypeDescription], [OrderTypeChannel], [AuditUpdateDate], [AuditUpdateBy])  
VALUES ('D ', 'Direct Sale', 'Direct', GETDATE(), SUSER_SNAME() )  
GO  
INSERT [LKP].[SalesInvoiceOrderType] (  
[OrderTypeCode], [OrderTypeDescription], [OrderTypeChannel], [AuditUpdateDate], [AuditUpdateBy])  
VALUES ('R ', 'Direct Credit', 'Direct', GETDATE(), SUSER_SNAME() )  
GO  
INSERT [LKP].[SalesInvoiceOrderType] (  
[OrderTypeCode], [OrderTypeDescription], [OrderTypeChannel], [AuditUpdateDate], [AuditUpdateBy])  
VALUES ('S ', 'Reload Sale', 'Reload', GETDATE(), SUSER_SNAME() )  
GO  
INSERT [LKP].[SalesInvoiceOrderType] (  
[OrderTypeCode], [OrderTypeDescription], [OrderTypeChannel], [AuditUpdateDate], [AuditUpdateBy])  
VALUES ('T ', 'Reload Direct', 'Reload', GETDATE(), SUSER_SNAME() )
```

Build action =
none since this is
DML

















Schema Compare

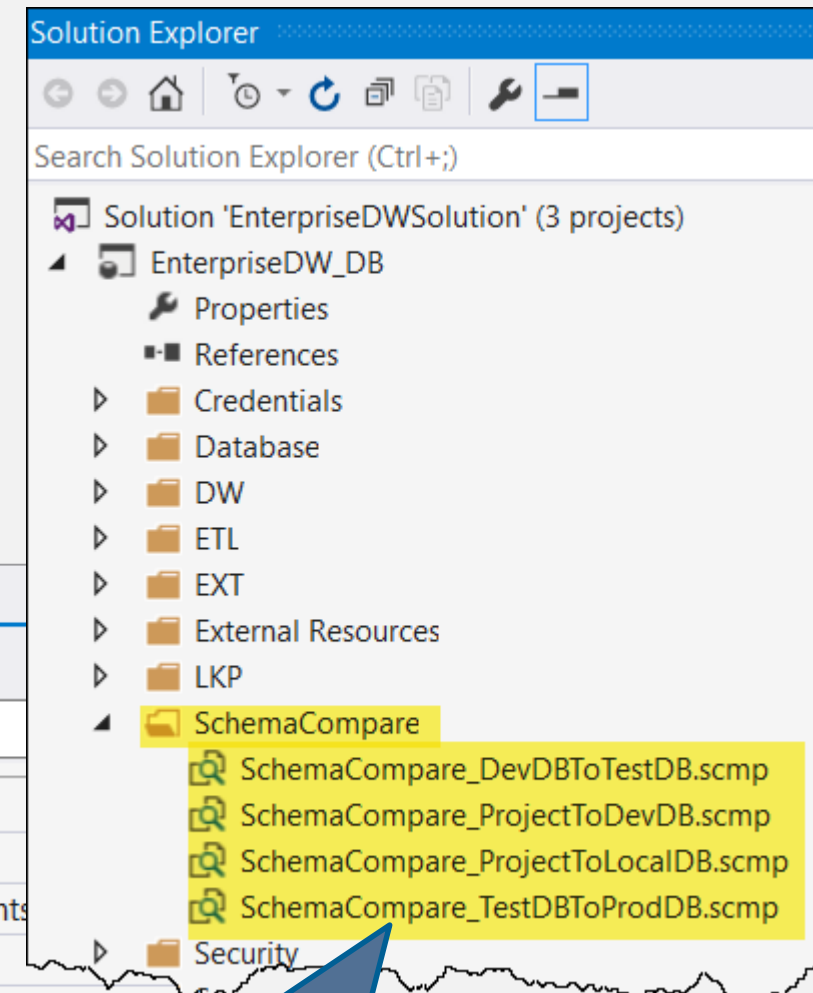
Settings to exclude permissions, users, etc + options to ignore

SchemaCompare_Pr...ctToLocalDB... x DimCustomer.sql [Design]

Compare Update     

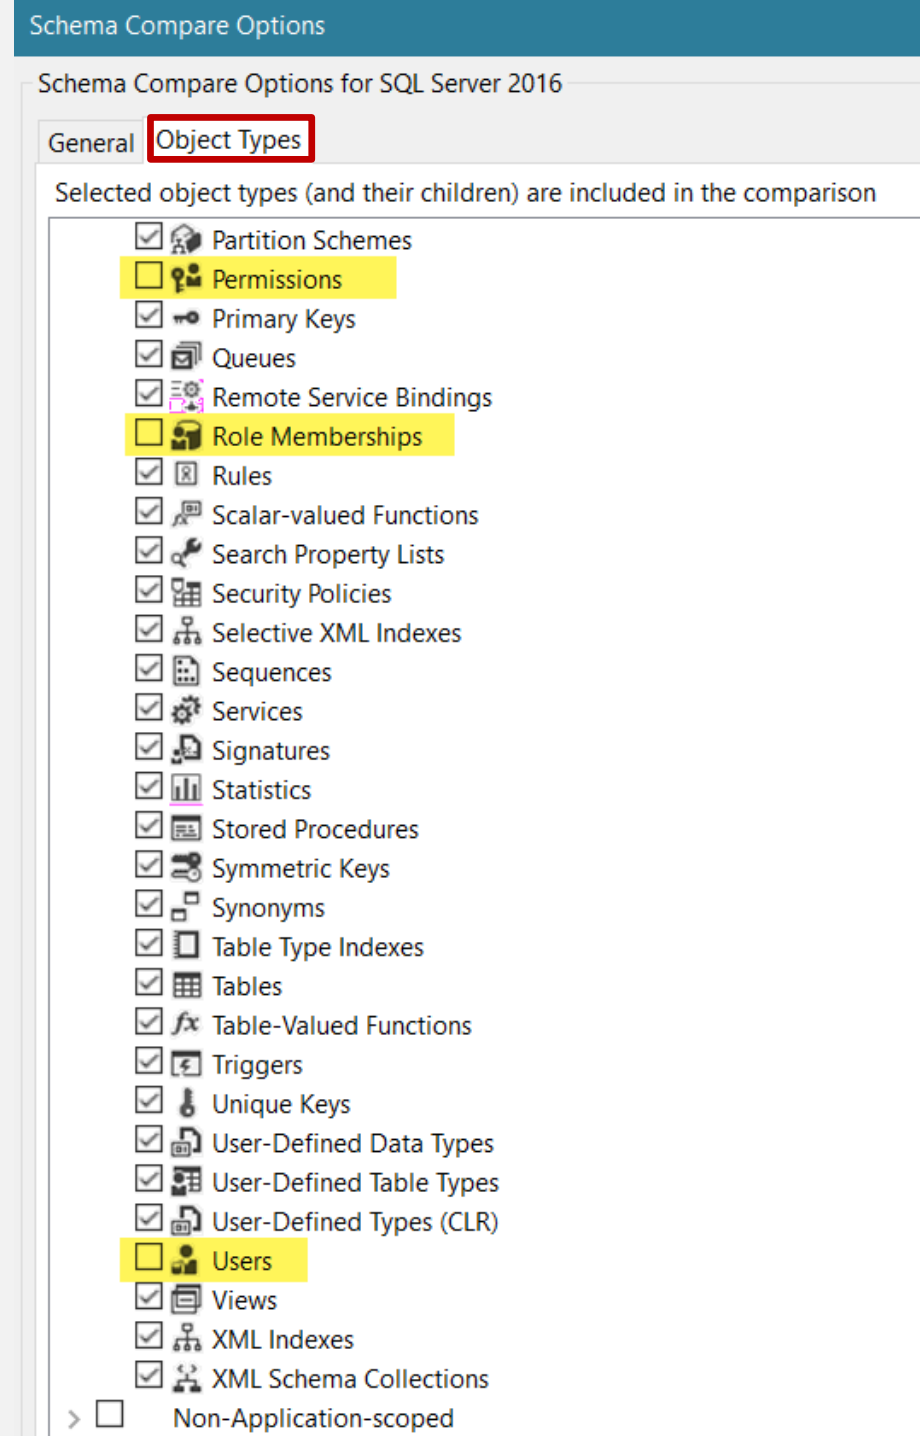
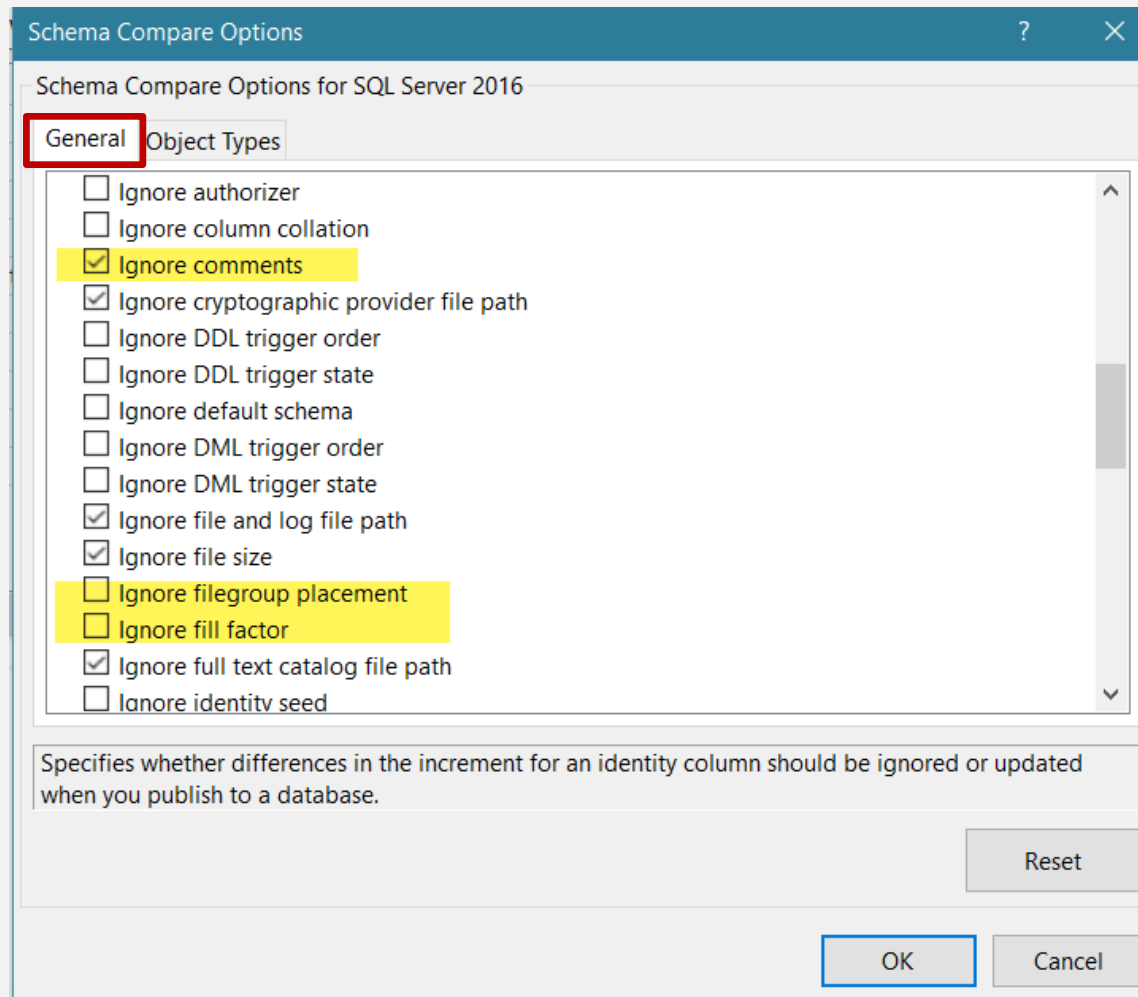
C:\Users\mcoates\OneDrive\Visual Studio Projects\De...\EnterpriseDW_DB ⇌ SSLAP40.EnterpriseDW

Type	Source Name	Action	Target Name
▼  Delete			
>  View		<input checked="" type="checkbox"/> 	USR.vwSensorMeasurements
>  External Data Source		<input checked="" type="checkbox"/> 	SensorDataDS
>  External File Format		<input checked="" type="checkbox"/> 	TextFileCommaDelimitedFF
>  External Table		<input checked="" type="checkbox"/> 	EXT.SensorMeasurements
▼  Change			
>  Table	DW.DimCustomer	<input checked="" type="checkbox"/> 	DW.DimCustomer
>  Table	DW.FactARSnapshot	<input checked="" type="checkbox"/> 	DW.FactARSnapshot
>  Table	DW.FactSalesInvoice	<input checked="" type="checkbox"/> 	DW.FactSalesInvoice



Saved settings

Schema Compare Options



Project Properties

Option to
generate error
during build

EnterpriseDW_DB x SqlSchemaCompare1* EnterpriseDW.sql

Project Settings
SQLCLR
SQLCLR Build
Build
SQLCMD Variables
Build Events
Debug
Reference Paths
Code Analysis

Configuration: Active (Debug) Platform: Active (Any CPU)

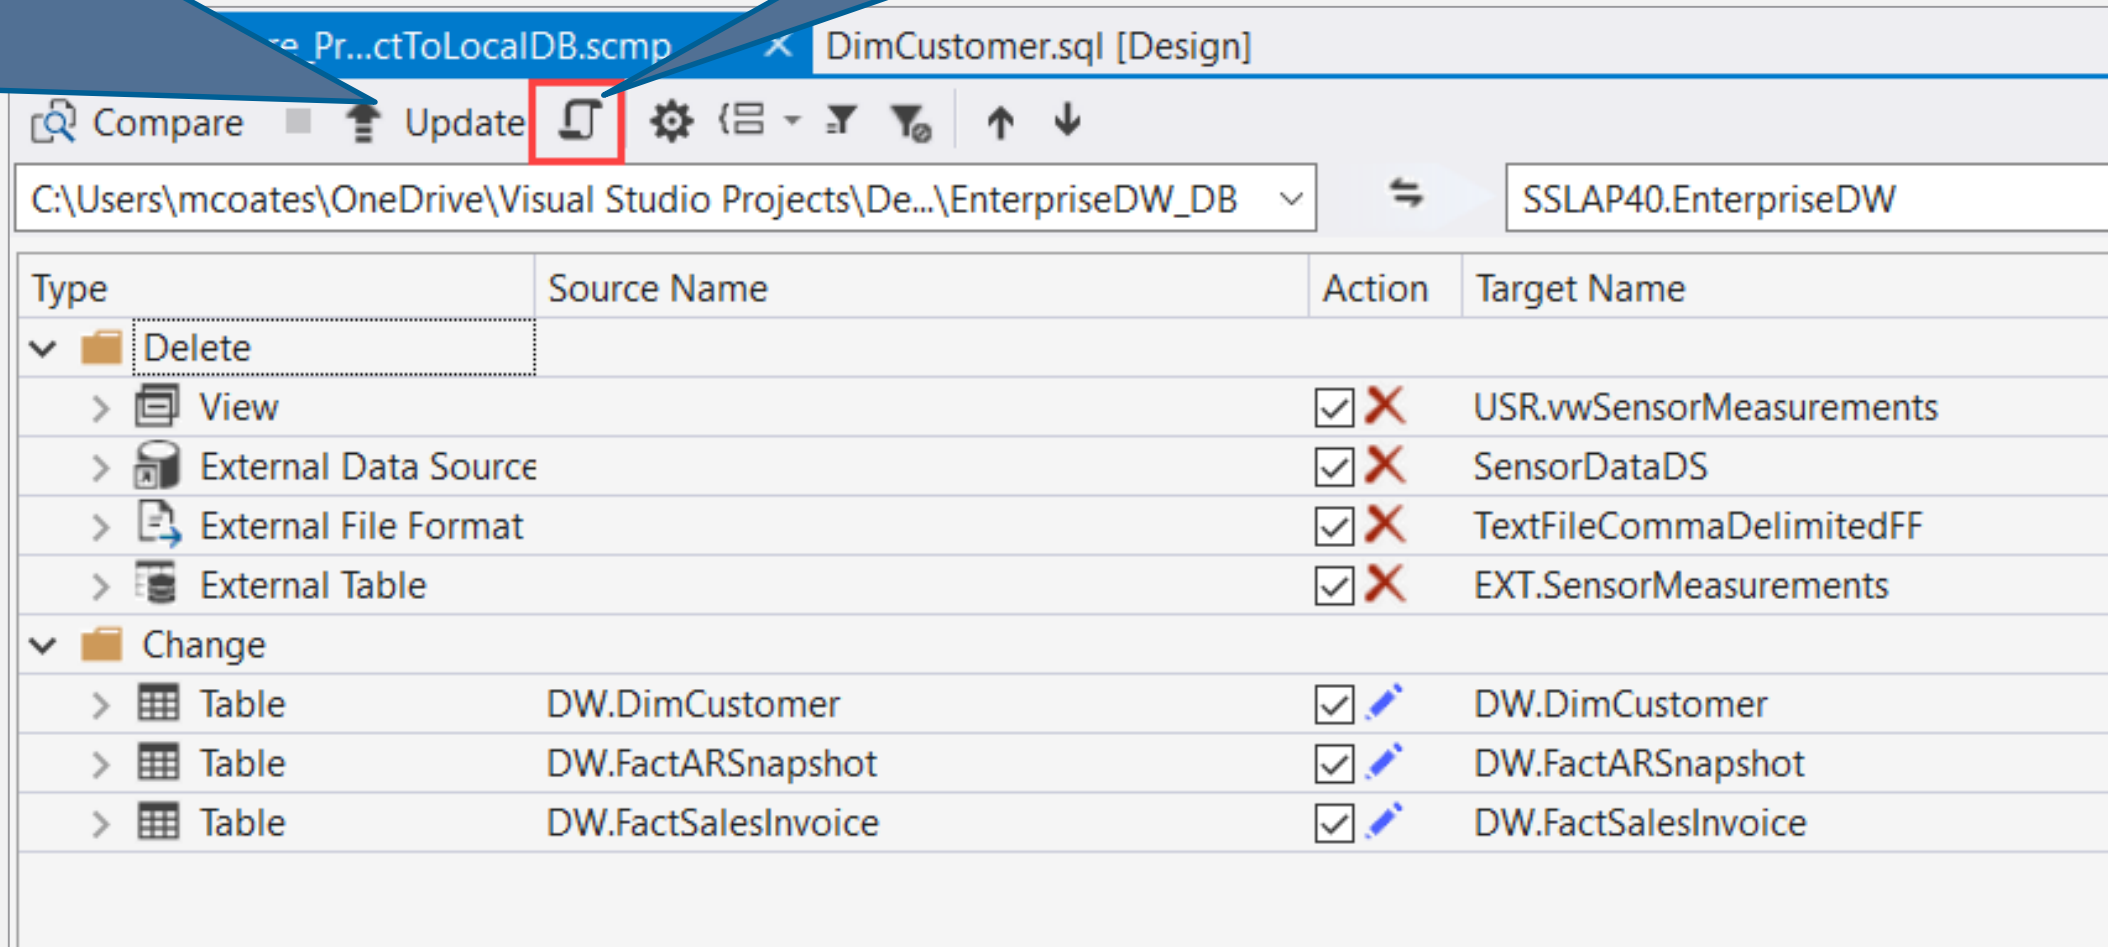
☐ Enable Code Analysis on Build








Rules	Treat Warning as Error
<input checked="" type="checkbox"/> Microsoft.Rules.Data.Design	<input checked="" type="checkbox"/>
<input checked="" type="checkbox"/> SR0001: Avoid SELECT * in stored procedures, views, and table-valued functions.	<input checked="" type="checkbox"/>
<input checked="" type="checkbox"/> SR0008: Consider using SCOPE_IDENTITY instead of @@IDENTITY.	<input type="checkbox"/>
<input checked="" type="checkbox"/> SR0009: Avoid using types of variable length that are size 1 or 2.	<input checked="" type="checkbox"/>
<input checked="" type="checkbox"/> SR0010: Avoid using deprecated syntax when you join tables or views.	<input checked="" type="checkbox"/>
<input checked="" type="checkbox"/> SR0013: Specify values for output parameters in all code paths.	<input type="checkbox"/>
<input checked="" type="checkbox"/> SR0014: Maintain compatibility between data types.	<input type="checkbox"/>
<input checked="" type="checkbox"/> Microsoft.Rules.Data.Naming	<input checked="" type="checkbox"/>
<input checked="" type="checkbox"/> SR0011: Avoid using special characters in object names.	<input checked="" type="checkbox"/>
<input checked="" type="checkbox"/> SR0012: Avoid using reserved words for type names.	<input checked="" type="checkbox"/>
<input checked="" type="checkbox"/> SR0016: Avoid using sp_ as a prefix for stored procedures.	<input type="checkbox"/>
<input checked="" type="checkbox"/> Microsoft.Rules.Data.Performance	<input type="checkbox"/>
<input checked="" type="checkbox"/> SR0004: Avoid using columns that do not have an index as test expressions in IN predicates.	<input type="checkbox"/>
<input checked="" type="checkbox"/> SR0005: Avoid using patterns that start with "%" in LIKE predicates.	<input type="checkbox"/>
<input checked="" type="checkbox"/> SR0006: In the comparison, simplify the expression that includes indexed columns.	<input type="checkbox"/>
<input checked="" type="checkbox"/> SR0007: Use ISNULL(column, default value) on nullable columns in expressions.	<input type="checkbox"/>
<input checked="" type="checkbox"/> SR0015: Extract deterministic function calls from WHERE predicates.	<input type="checkbox"/>

Schema Compare


Usually don't want to let the target update directly

Generates a script to use for deployment

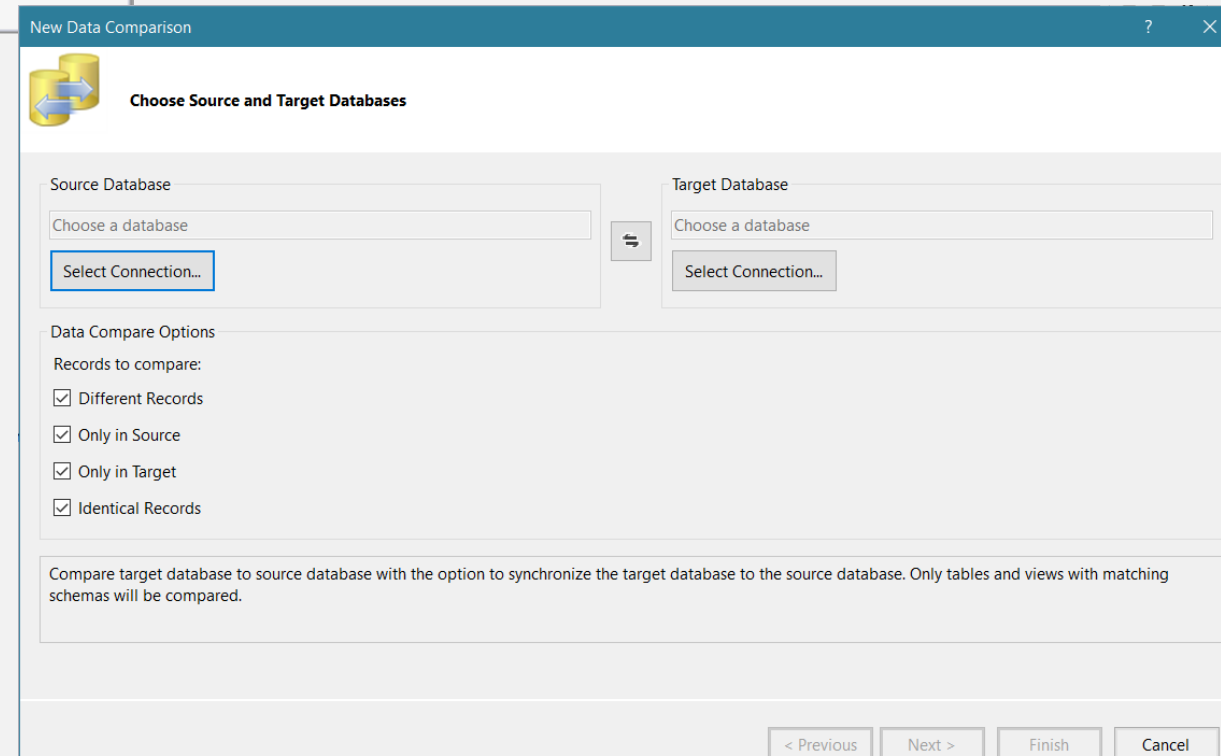


Type	Source Name	Action	Target Name
▼ Delete			
> View		<input checked="" type="checkbox"/> 	USR.vwSensorMeasurements
> External Data Source		<input checked="" type="checkbox"/> 	SensorDataDS
> External File Format		<input checked="" type="checkbox"/> 	TextFileCommaDelimitedFF
> External Table		<input checked="" type="checkbox"/> 	EXT.SensorMeasurements
▼ Change			
> Table	DW.DimCustomer	<input checked="" type="checkbox"/> 	DW.DimCustomer
> Table	DW.FactARSnapshot	<input checked="" type="checkbox"/> 	DW.FactARSnapshot
> Table	DW.FactSalesInvoice	<input checked="" type="checkbox"/> 	DW.FactSalesInvoice

Data Compare



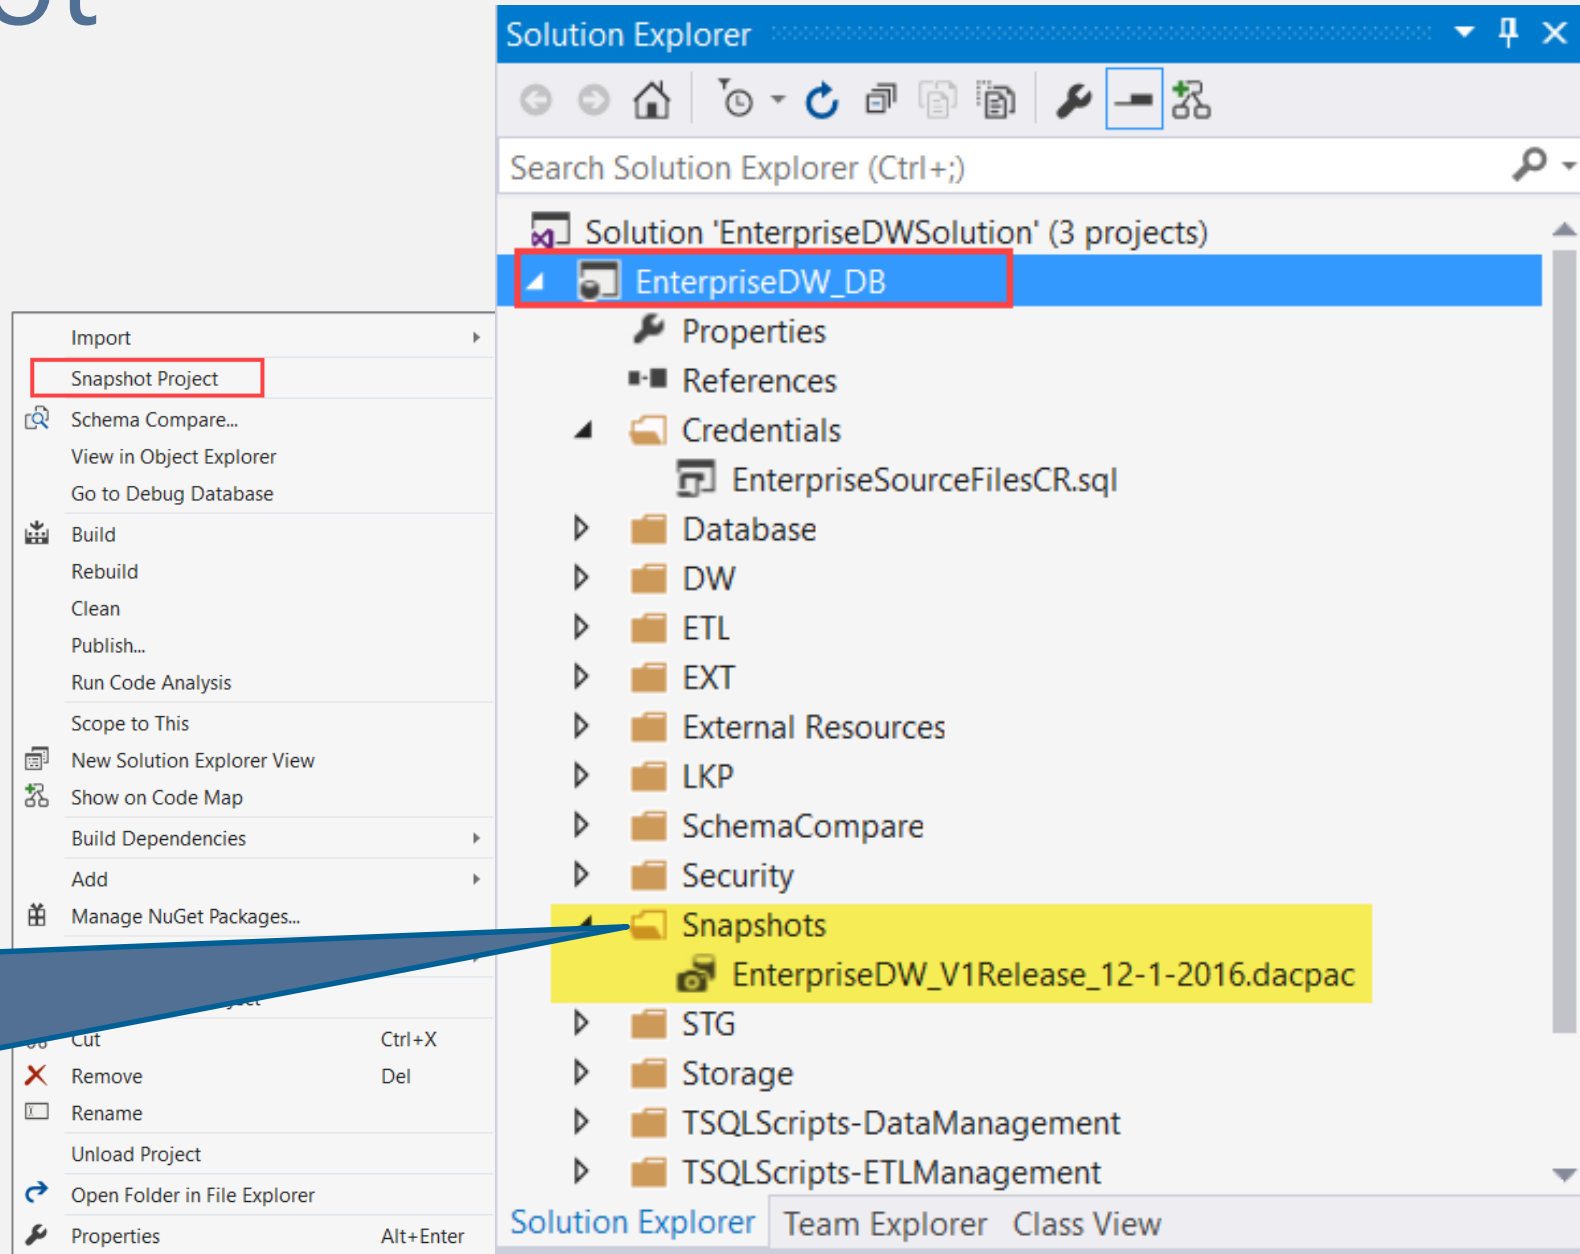
Basic functionality to compare data between two tables -- schema must match.



Project Snapshot

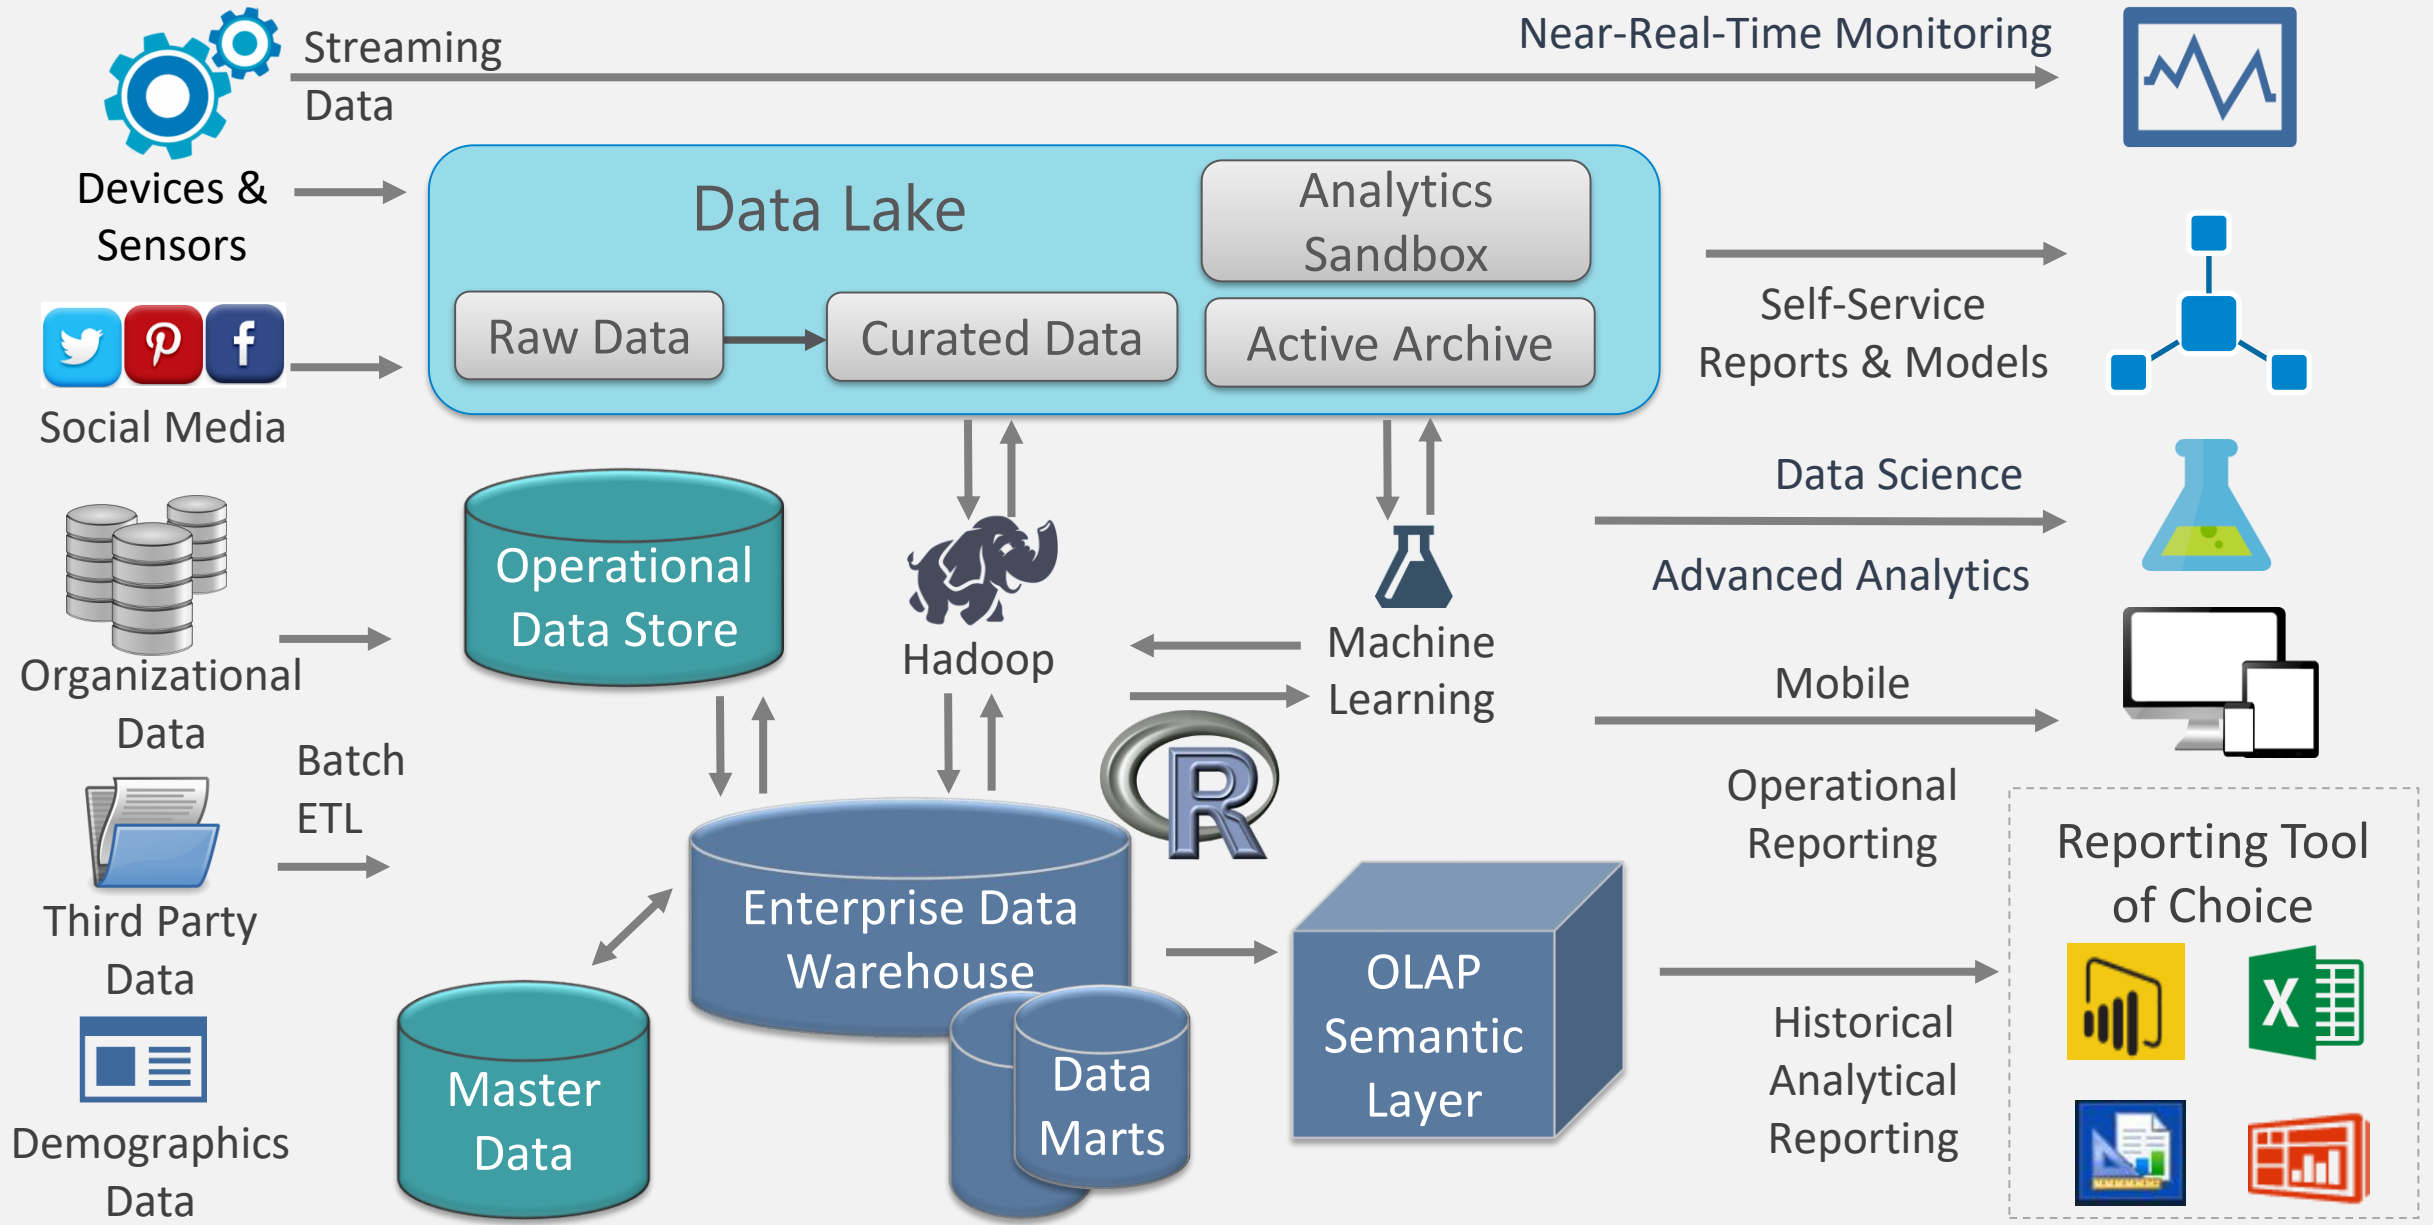
Snapshot of the database schema at a point in time (ex: major release points).

Store the
.dacpac file in
the project if
desired



Planning Future Growth of the Data Warehouse

Modern /DW/BI/Analytics Systems



Growing your DW/BI/Analytics Environment



Cloud &
Hybrid
Platforms



Modern DW
Multi-Platform
Architecture



Advanced
Analytics



Real-Time
Reporting



Self-
Service
BI



Agile,
Nimble
Solutions

Achieving Extensibility in a DW

Design with change in mind. Ex: Create a lookup table with code/descriptions, or implement in a view, rather than hard-coding in ETL.

Plan for a **hybrid** environment with **multiple architectures**.

Introduce **conformed dimensions first** whenever possible.

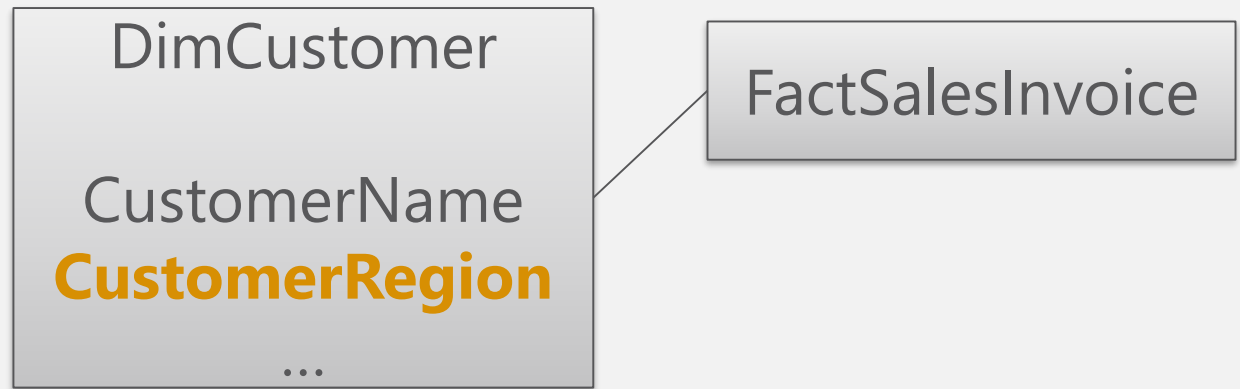
Try to **avoid isolated “stovepipe” implementations** unless the isolation is absolutely intended.

Conduct **active prototyping sessions** with business users to flush out requirements. A data modeling tool like Power BI works well for this.

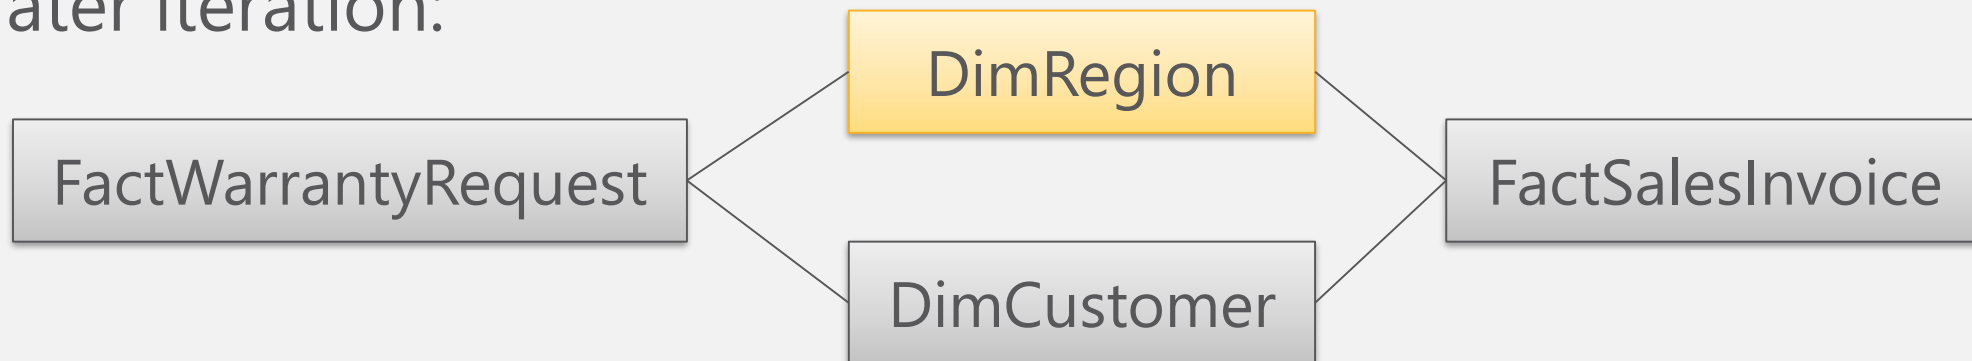
Achieving Extensibility in a DW

Be prepared to do some **refactoring** along the way. Ex: converting an attribute to be a conformed dimension.

First implementation:



Updated in a later iteration:



Achieving Extensibility in a DW

Introducing new measures:

- Can be a new column in a fact table as long as it's the same grain & the same business process

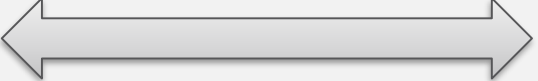
Introducing new attributes:

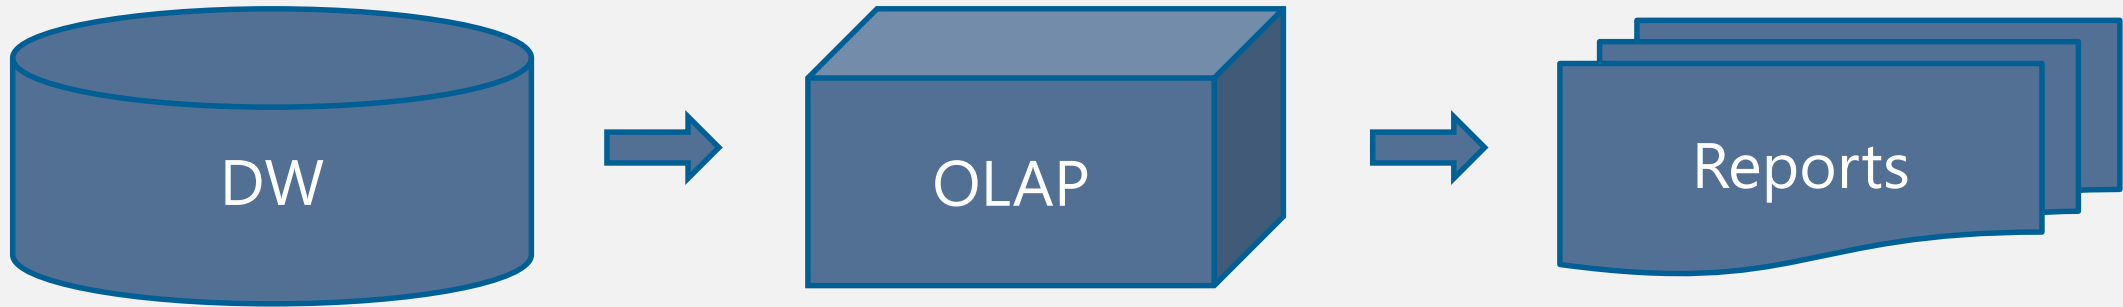
- Can be a new column in a dimension, or
- Can be via a new foreign key in a fact table as long as it doesn't affect the grain

Agility for the things that usually require the most time investment:

- Data modeling
- ETL processes
- Data quality

Achieving Extensibility in a DW

Reusability Downstream  Speed of Change Implemented



Consider using an ***OLAP cube or in-memory model*** (like Analysis Services) for:

- Summary data (as opposed to summary tables in your DW)
- Year-to-Date type of calculations
- Year-over-Year type of calculations
- Aggregate level calculations (as opposed to row-by-row calculations)

Modern DW: Important Concepts to Know

Polygot Persistence

Using the most effective data storage technology to handle different data storage needs

Lambda Architecture


Data processing architecture which supports large amounts of data via a speed layer, batch layer, and serving layer

Schema on Read





Data structure is applied at query time rather than when the data is initially stored

Recommended Resources

Read
First




Read
Second



Thank You for Attending

To download a copy of this presentation:
SQLChick.com "Presentations & Downloads" page



Melissa Coates
BI Architect, SentryOne
sentryone.com

Blog: sqlchick.com
Twitter: @sqlchick

*Creative Commons License:
Attribution-NonCommercial-NoDerivative Works 3.0*

